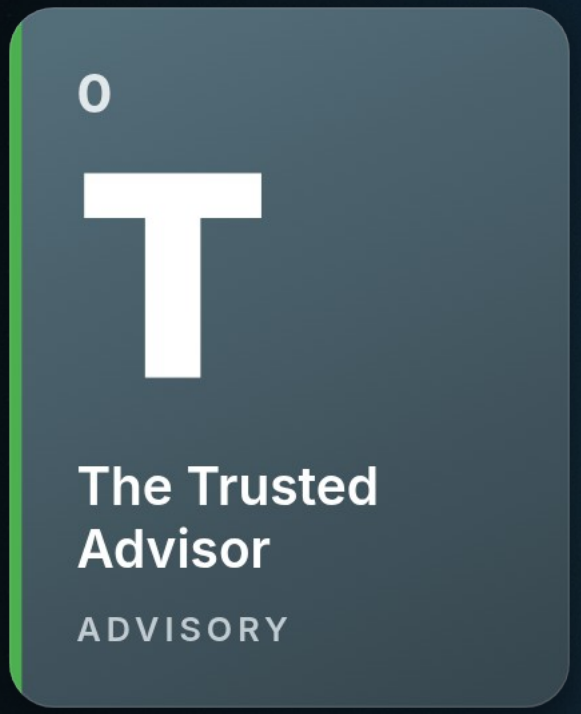

TICINO DATA CONFERENCE 2026 · BIOGGIO

Boss Your Agents: How AI Secures What AI Teams Build

Matthias Falland

MICROSOFT DATA PLATFORM MVP · THE TRUSTED ADVISOR

ABOUT ME

Matthias Falland

The Trusted Advisor · Microsoft Data Platform MVP since 2018

- **Microsoft MVP** Data Platform, since 2018
- **Fabric Friday** weekly on YouTube, one Fabric element per episode
- **Fabric Architecture Podcast** weekly, with an AI co-host
- **Fabric Periodic Table** fabricperiodictable.com, every Fabric item
- **Meetups** Zurich, Hamburg, Basel: over 5,000 members

LinkedIn: matthias-falland
Repo of this talk →

What an agent session is

PROGRAM · CLAUDE CODE IN A TERMINAL · ONE SESSION, ONE RUN

00 Why rules go missing

context, not intelligence

- The model follows what it reads: your rules are in a wiki and in people's heads
- Without them: missing row-level security, inconsistent names, undocumented decisions
- The five steps put the project's knowledge into the context, and check it

THE ARC

Zero to hero, in five steps

THE MAIN CHARACTER IS THE PROJECT'S KNOWLEDGE

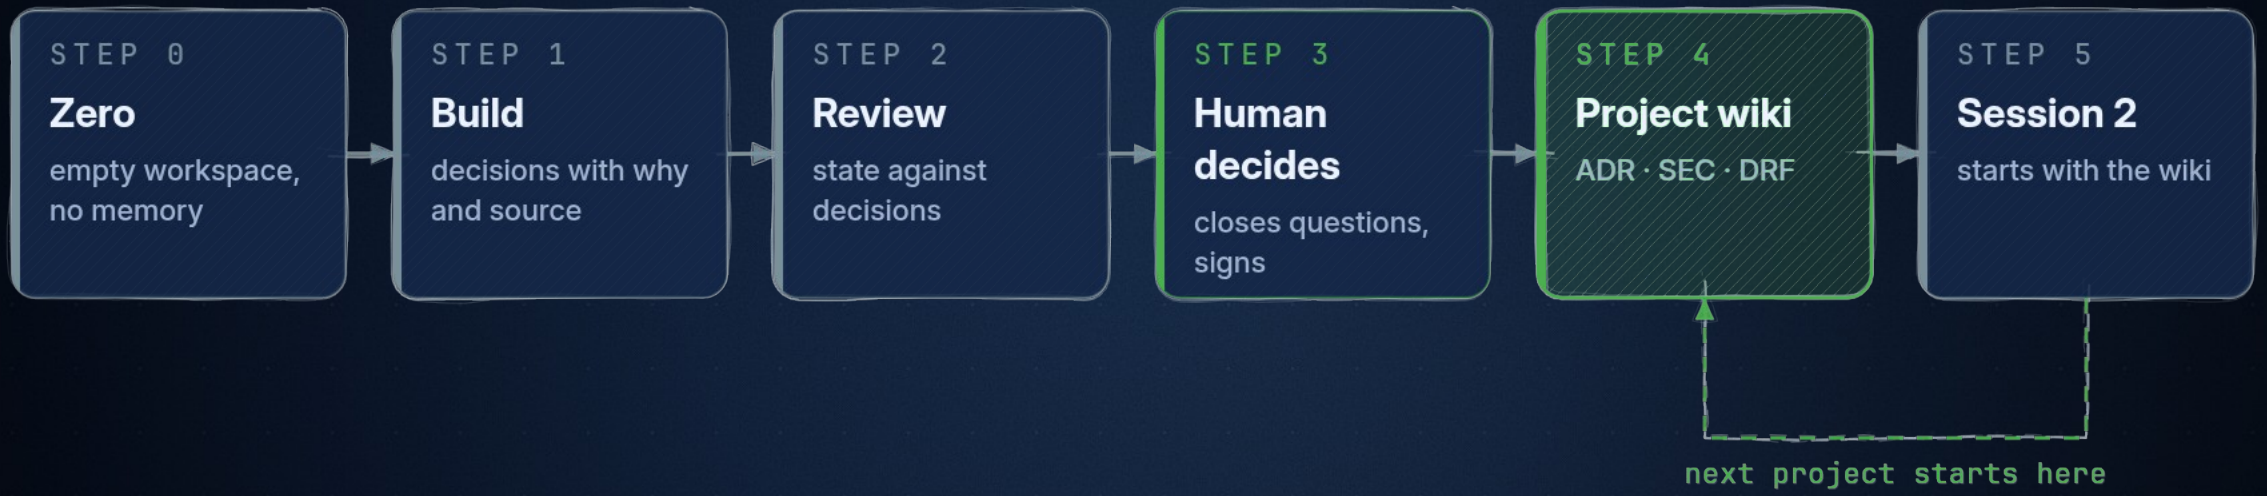

POSITIVE PATH · NOTHING SEEDED · EVERY STEP IS A FILE OR A COMMAND YOU CAN REDO

Zero

What does a session need before it builds anything?

37

WS

Workspace

GOVERN

STEP 0 · MECHANISM

Three layers, set before the first prompt

	IDENTITY	WORKSPACE ROLE	TOOLS AND MCP SCOPE
Builder session	**Its own account** never yours, never an admin	**Contributor, one workspace** creates no workspaces	**MS Learn + wiki + toolkit** knowledge, memory, Fabric
Reviewers	**Their own account** separate from the builder	**Contributor, to read definitions** the session keeps it from writing	**Read checks + wiki** no build tools
Human	**You** the only one who decides	**Approves and signs** the baseline	**Reads the report** answers open questions

FACT CARD · STEP 0

Zero

WHAT

A session that starts empty: its own profile, its own Fabric identity, two tools. Microsoft Learn for knowledge, a wiki for memory.

GOOD WHEN

You start a project, or hand an existing one to an agent for the first time.

FILES & COMMANDS

```
.mcp.json
  microsoft-learn + wiki
CLAUDE.md
  record every decision:
  what, why, alternatives, source
./run.sh
  clean profile, own identity
```

TAKEAWAY **A decision log needs an instruction and a place to write.**

1 Build

How does an agent leave its decisions behind?

STEP 1 · MECHANISM

Every decision becomes a page

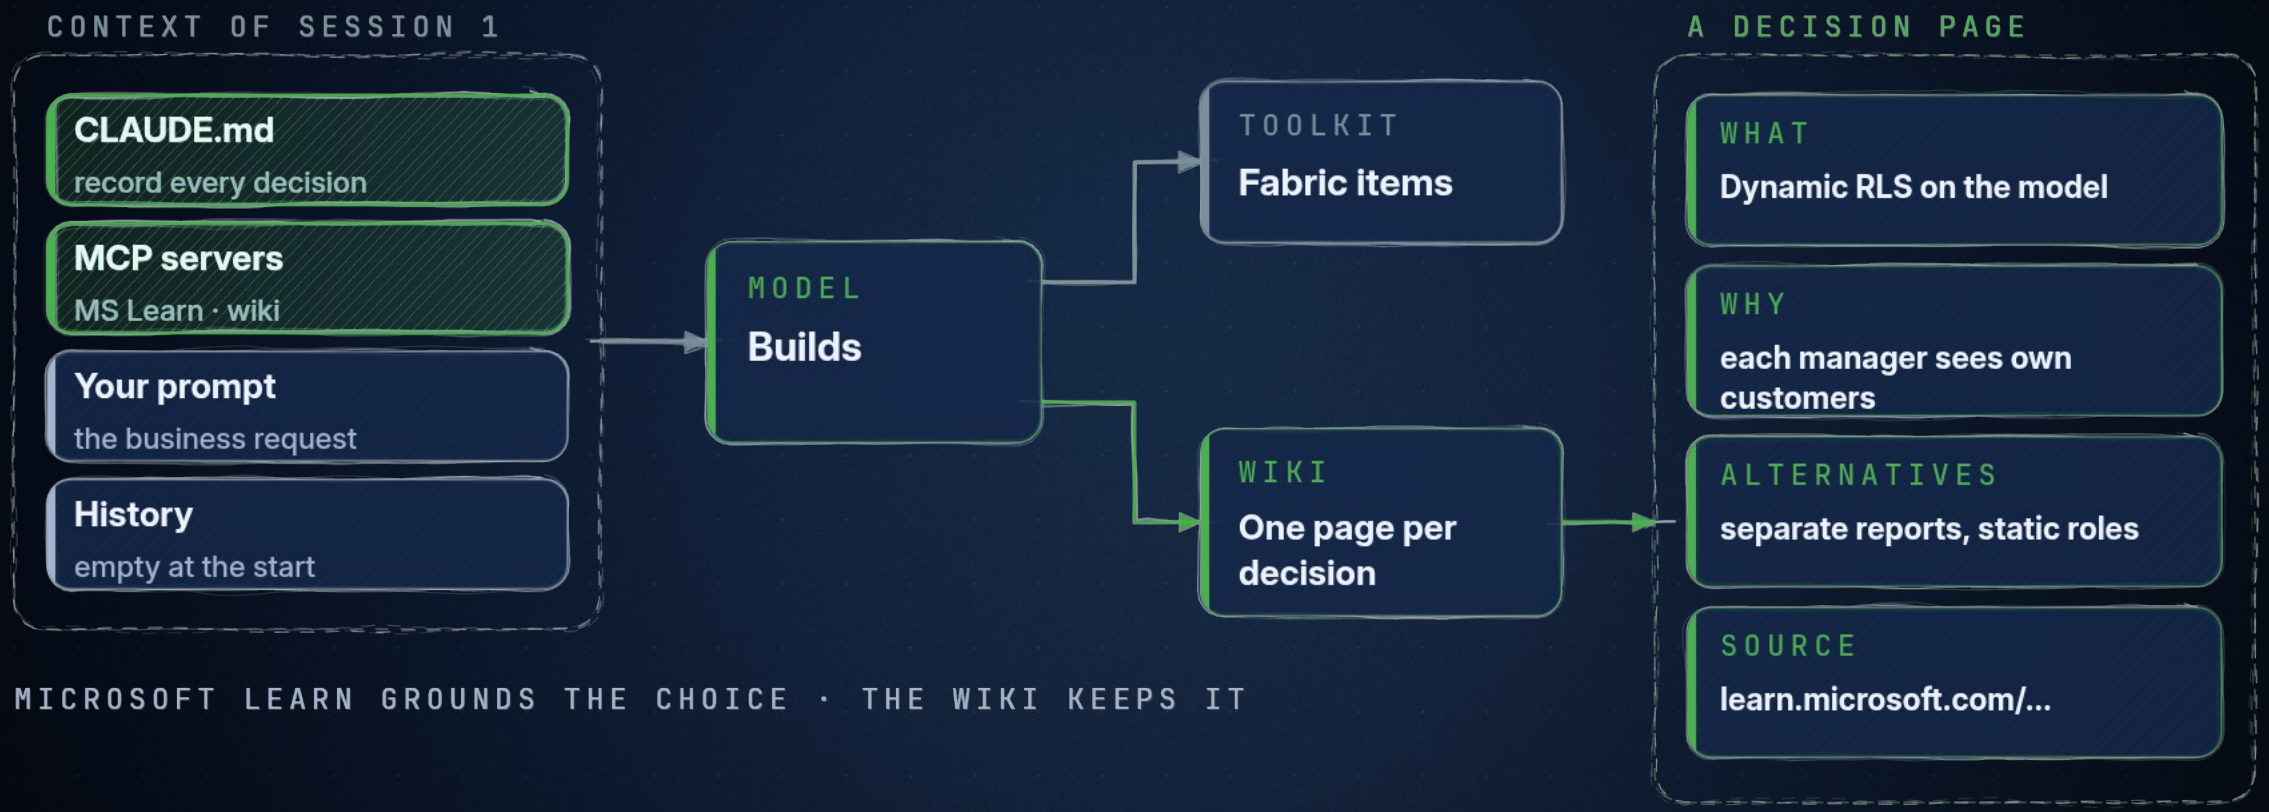

MICROSOFT LEARN GROUNDS THE CHOICE · THE WIKI KEEPS IT

Step 1 · Session 1 builds and records its decisions

```
index.md
lakehouse-sales-lh-no-schemas-delta-written-over-onelake.md
log.md
sales-model-access-role-membership-done-sharing-and-fixed-identi.md
sales-model-shape-star-with-two-tables-explicit-measures.md
$ glow -p "$(ls .omc/wiki/*{rls,row-level}* 2>/dev/null | head -1)"
$ glow -p "$(ls .omc/wiki/*direct-lake* | head -1)"
```

Direct Lake on OneLake for the Sales model

Decision: The Sales semantic model uses Direct Lake on OneLake (shared expression `AzureStorage.DataLake` pointing at the lakehouse), not Direct Lake on SQL analytics endpoint.

Why: Microsoft recommends Direct Lake on OneLake for new semantic models. It runs only in Direct Lake mode, so it never falls back to DirectQuery, and its query plans are more efficient. We don't need any SQL-endpoint security or SQL views, which are the only reasons to pick Direct Lake on SQL.

Alternatives considered: Direct Lake on SQL (`Sql.Database`) can fall back to DirectQuery and adds a dependency on the SQL endpoint. Import mode doesn't meet the request, which asks for Direct Lake.

Source: **https://learn.microsoft.com/fabric/fundamentals/direct-lake-develop** https://learn.microsoft.com/fabric/fundamentals/direct-lake-develop (Choose the right storage mode); **https://learn.microsoft.com/fabric/fundamentals/direct-lake-how-it-works#directquery-fallback** https://learn.microsoft.com/fabric/fundamentals/direct-lake-how-it-works#directquery-fallback

```
$
```

④ **Takeaway**

Read the pages yourself. A good page says what was decided, why, what else was possible, and where it is documented.

FACT CARD · STEP 1

Build with a decision log

WHAT

The builder works on the request and writes one wiki page per decision while it builds.

GOOD WHEN

An agent makes a choice that someone will ask about later.

FILES & COMMANDS

```
/build
  read task/, deliver it
.omc/wiki/
  one page per decision:
  what · why ·
  alternatives · source
```

TAKEAWAY **A good page says what, why, what else was possible, and the source.**

2
Review

Does the state match the decisions, and what was never decided?

STEP 2 · MECHANISM

What the reviewer gets

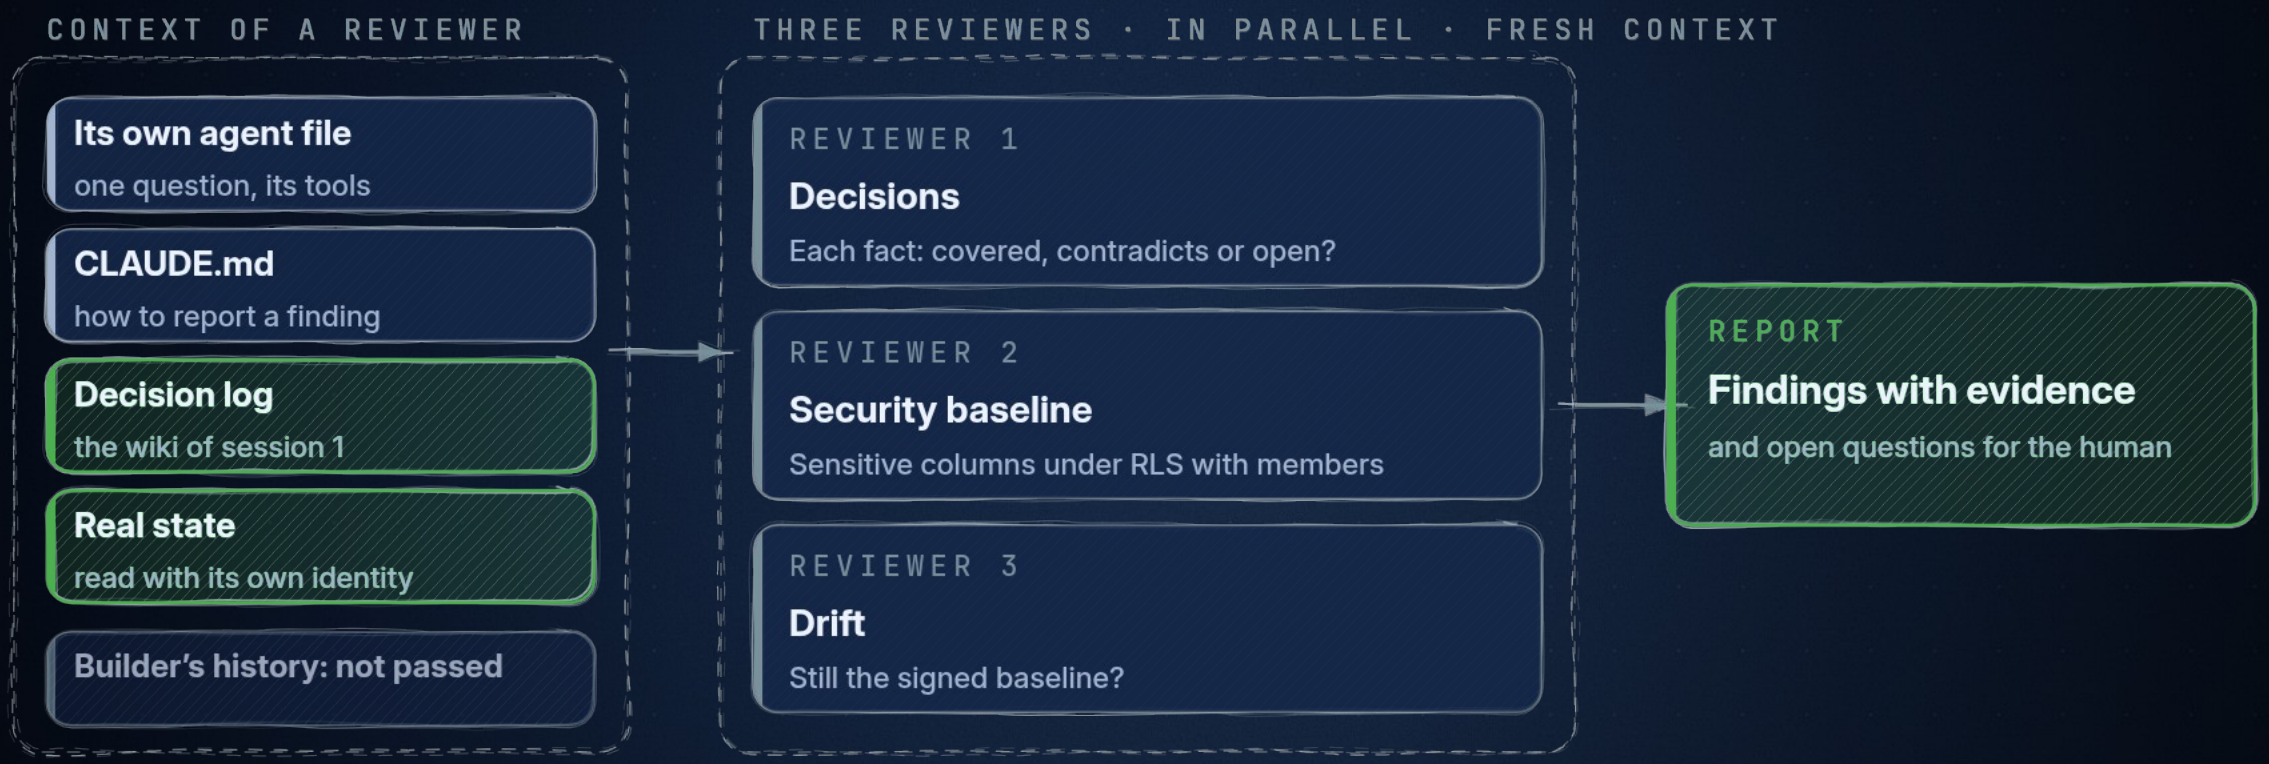

STEP 2 · EVIDENCE

A finding is a claim, evidence and a remedy

CLAIM

What the reviewer says

one sentence, checkable

EVIDENCE

What was read, by whom

item, definition, wiki page · reviewer identity

REMEDIATION

What would close it

a change, or a decision

EXAMPLE FROM THE REHEARSALS

OPEN QUESTION

Item names were never decided.

EVIDENCE

Read: the workspace item list and every page of the decision log

as the reviewer's own identity · no page mentions naming

REMEDIATION

A human decides the convention; it goes into the wiki.

A GAP IS A QUESTION FOR THE HUMAN, NOT A TRAP

Step 2 · Three reviewers read the real state

What naming convention should items follow (sales_lh, Sales), and why these names?

• Evidence: F1: items built: Lakehouse «sales_lh»; SemanticModel «Sales». Pages use the names but none decides naming. No decision page covers it.
• Remediation: a human decides (one sentence) and the decision goes into the log
• Reviewer: decisions

Info

SEC-010 · baseline · FF Ticino Session 1 DEV

workspace roles not readable by the reviewer (least privilege)

• Evidence: GET roleAssignments as **reviewer@contoso.example** mailto:reviewer@contoso.example: HTTP 403 InsufficientWorkspaceRole. RLS does not apply to Admin, Member or Contributor; a human checks who holds those roles.
• Reviewer: security

DRF-000 · - · FF Ticino Session 1 DEV

no baseline yet

• Evidence: no signed snapshot for this workspace; 2 item(s) read, one is signed after human approval
• Reviewer: drift

Next: a human answers the open questions (one sentence each) and approves; then `./py gate/snapshot.py sign` records the approved state.

$

④ Takeaway

Read the gate report before anyone builds on the result. Ask the reviewers for gaps as well as mistakes.

FACT CARD · STEP 2

Review with a fresh context

WHAT

Reviewers with their own identity and a fresh context read the deployed state and the decision log.

GOOD WHEN

Before anyone builds on a result, and after every change.

FILES & COMMANDS

```
.claude/agents/
  security · decision · drift
/review
  three in parallel, one gate
reports/latest-gate.md
  findings, open questions Q-n
```

TAKEAWAY

Ask for gaps, not for mistakes: a missing decision is an open question.

3
Human decides

Who turns a decision into a rule?

STEP 3 · MECHANISM

Close the question in your own words, then sign

SIGN-OFF

OPEN QUESTION

Item names were never decided

DECIDED, IN MY WORDS

Items use type prefixes: lh_, sm_, rpt_, nb_

WHY

Anyone can tell an item's type from its name

SIGNED

the baseline: state and decisions, today

THE AGENT

Proposes and records

decisions with why, alternatives and source

THE HUMAN

Decides what becomes a rule

closes open questions, signs the baseline

THE SIGNATURE

A fixed point to compare with

any later change shows up as drift

Step 3 · The open question closed, the baseline signed

```
$ cd ../review
$ ./py gate/snapshot.py sign
snapshot signed by Matthias Falland: 2 items (Lakehouse/sales_lh, SemanticModel/Sales),
 8 definition parts -> /home/demo/.local/state/zero-to-hero/·····················
·············/snapshot.json, reports/baseline.json
$
```

④ Takeaway

Agents propose and record. A person decides what becomes a rule and signs the state.

FACT CARD · STEP 3

The human decides

WHAT

The human reads the report, closes each open question in own words and signs the baseline.

GOOD WHEN

A choice should become a rule for the next project.

FILES & COMMANDS

```
/decide <one sentence>
  answers Q-1 as a wiki page,
  decided by you
./py gate/snapshot.py sign
  only after PASSED,
  records who approved
```

TAKEAWAY

Agents propose and record. You decide what becomes a rule.

4
Project wiki

Where does the project's knowledge live?

STEP 4 · MECHANISM

One folder, four kinds of pages

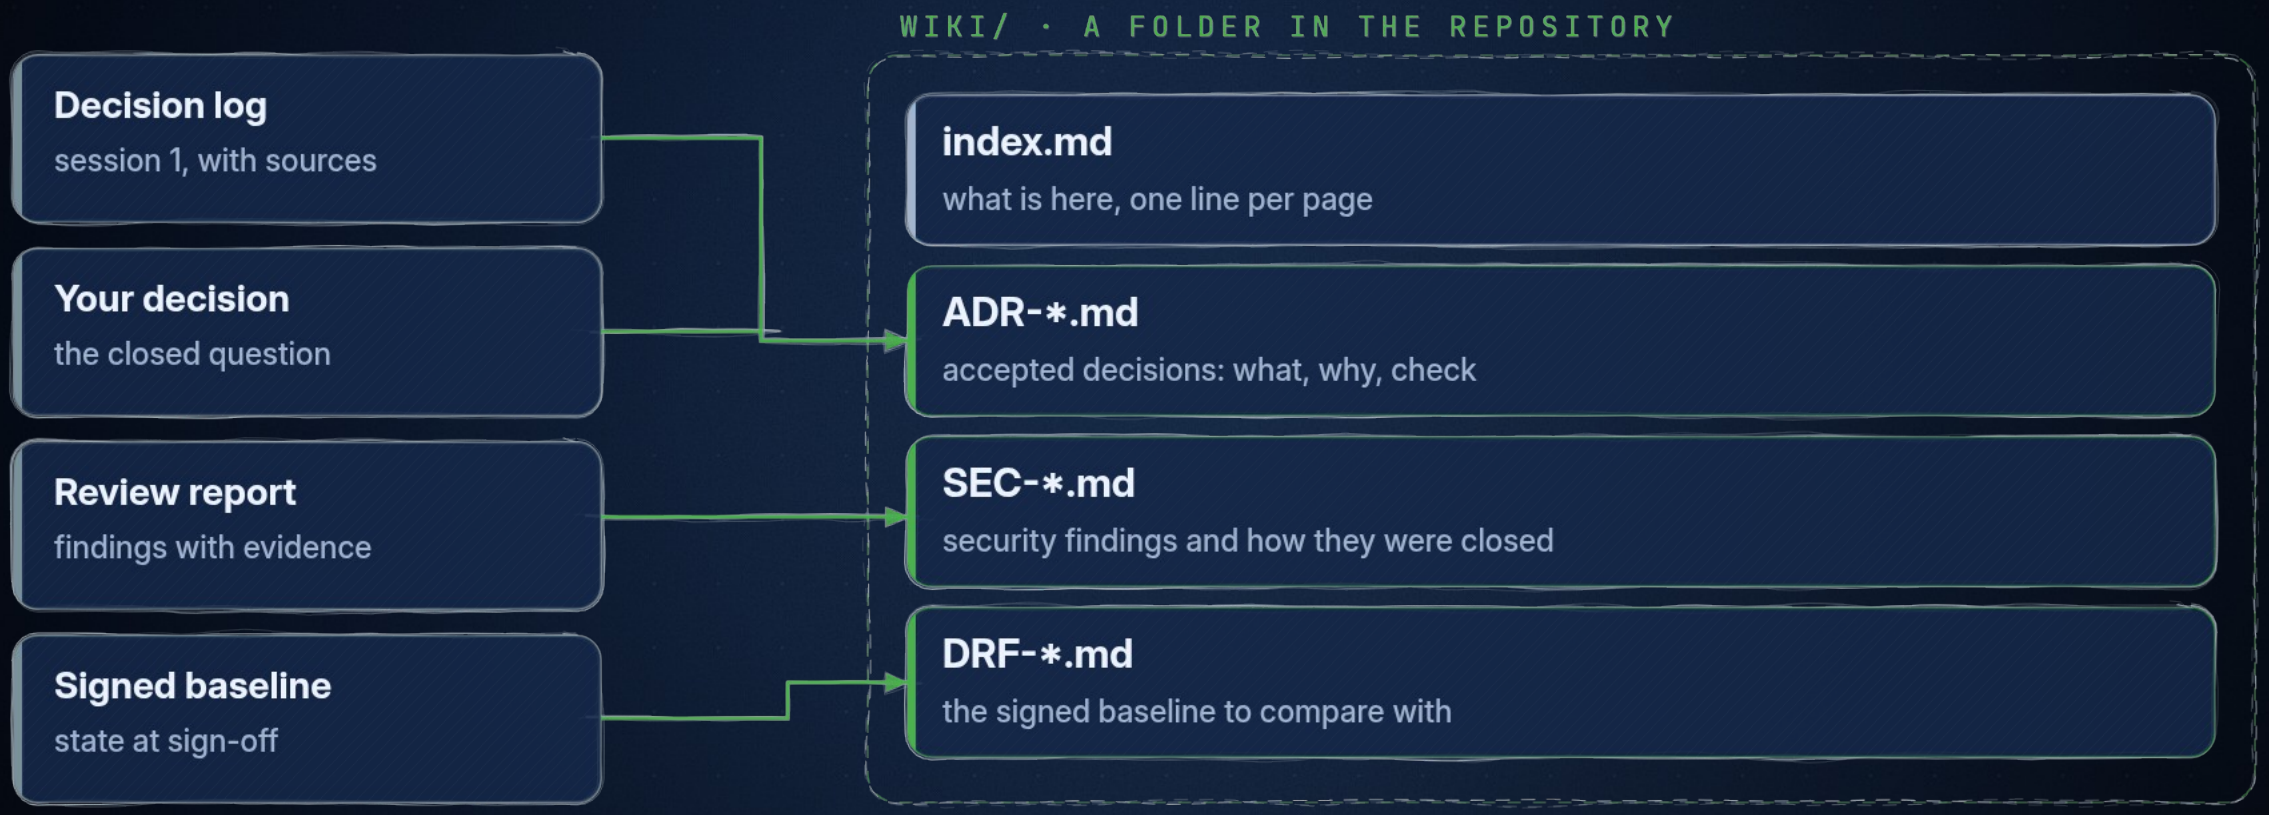

Step 4 · From log, report and decision to the wiki

ADR-005 Item naming with type prefixes

Status

Accepted. Decided by a human: Matthias Falland, project owner, 2026-10-27. Answers
open question Q-1 of the gate review of 2026-10-27 12:08.

Context

The review found the deployed items named without a common convention, and no decision page covered naming (Q-1): Lakehouse «sales_lh» carries its type as a suffix; SemanticModel «Sales» has no type marker and is capitalised.

Decision

"Items use type prefixes: lh_, sm_, rpt_, nb_, lower case, underscores, so anyone can tell an item's type from its name." Lakehouse `lh_`, SemanticModel `sm_`, Report `rpt_`, Notebook `nb_`.

The decision applies from the next delivery; nothing in the workspace changes now. Two deployed items do not follow it yet: «sales_lh» (suffix, not the prefix `lh_`) and «Sales» (no `sm_`, not lower case).

Why

So anyone can tell an item's type from its name (the project owner's reason). The convention stays within Fabric's naming rule: names begin with a letter and contain only alphanumeric characters and underscores.

:

④ Takeaway

A rule belongs in the wiki when a reviewer can check it. Everything else stays a recorded decision.

FACT CARD · STEP 4

The project wiki

WHAT

Decision log, review and your decisions become one wiki: accepted decisions, security findings, the signed baseline.

GOOD WHEN

A project ends, or the next team or agent starts.

FILES & COMMANDS

```
/compile-wiki
  ADR-* · SEC-* · DRF-*
  + rules/*.yaml
./publish-wiki.sh
  .omc/wiki → wiki/
wiki/index.md
  one line per page
```

TAKEAWAY

A rule belongs in the wiki when a reviewer can check it.

5
Session 2

What changes when the next session starts with the wiki?

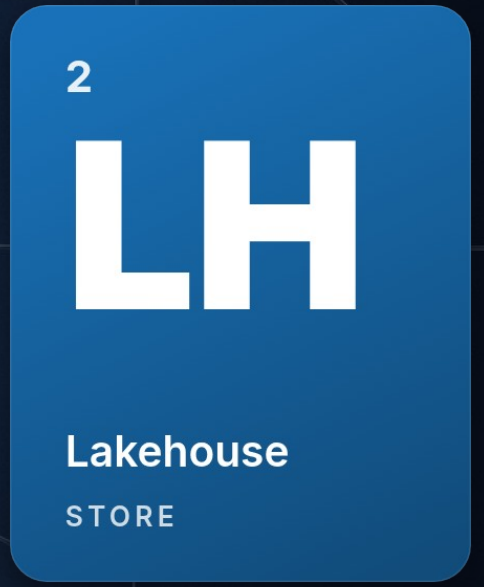

STEP 5 · MECHANISM

The next session starts with the wiki

Step 5 - Session 2 plans with the wiki, the review confirms

Plan: Sales Operations request (session 2)

Request: `task/sales-request.md`. Workspace: FF Ticino Session 2 DEV (empty at start). Builder identity:
ff-demo-contributor. Follows the accepted decisions of the last delivery: ADR-001..ADR-005, SEC-010, DRF-000.

Items

Item	Type	Name	Why
Lakehouse	Lakehouse	`lh_sales`	ADR-005 type prefix `lh_` …
Delta tables	in `lh_sales`	`customers` (60 rows), …	ADR-002
Semantic mod…	SemanticModel	`sm_sales` (TMDL in `wor…`	ADR-005 prefix `sm_` ; ADR…

Steps

1. `lakehouse create lh_sales` without schemas (ADR-002, rule `lakehouse_schemas: false`).
2. Upload `data/customers.csv`, `data/orders.csv` to `Files/raw/`; `load-table` with the default
`--engine onelake` (ADR-002: Spark Load Table API throttles the F2 with HTTP 430).
3. Model `sm_sales`, Direct Lake on OneLake: shared expression `AzureStorage.DataLake` on
`https://onelake.dfs.fabric.microsoft.com/<workspace id>/<lakehouse id>` (ADR-001), compatibility level 1604.
4. Shape (ADR-003): star, dimension `customers` (Customer Name, Region, Account Manager, Credit Limit; key hidden),
fact `orders` (Order Date, hidden amount and keys), single-direction many-to-one
:

④ Takeaway

The plan in `work/PLAN.md` cites the ADR ids before anything is built. Hand the knowledge on, and keep the review.

FACT CARD · STEP 5

Session 2 starts with knowledge

WHAT

A new session loads the wiki, plans with references to the ADRs, builds, and the same review confirms.

GOOD WHEN

Every new project, and every new agent on an old one.

FILES & COMMANDS

```
./handover.sh
  wiki/ + one line in CLAUDE.md:
  @wiki/index.md
/governed-build
  plan cites ADR ids, builds,
  review against wiki/rules/
```

TAKEAWAY

Hand the knowledge on, and keep the review.

FOUR PATTERNS

Where the four patterns sit in the five steps

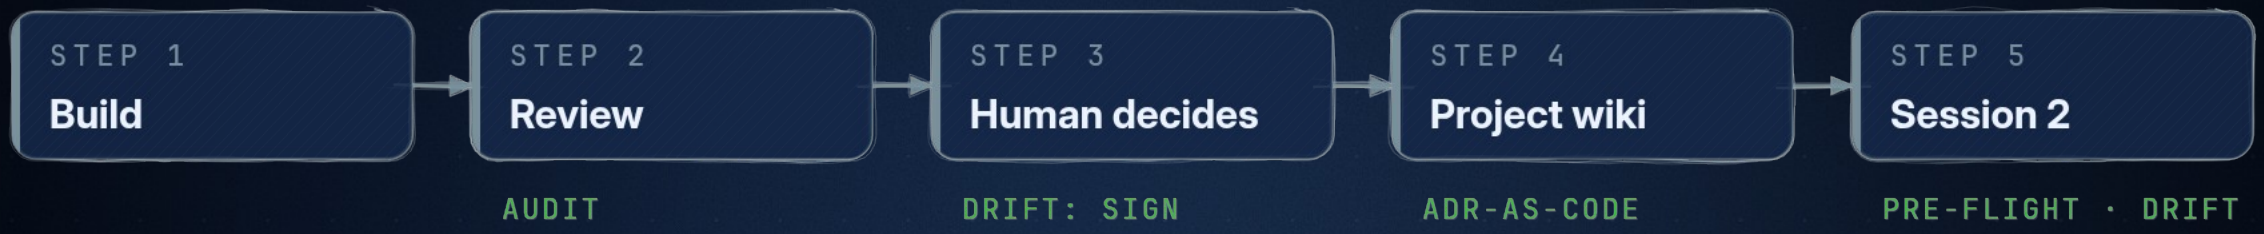

STEP 2

Post-deployment audit

Reviewers read what is deployed, not what the builder reports.

STEPS 3 AND 5

Drift detection

Signed baseline at step 3, compared again at step 5 and after any hand change.

STEP 4

ADR-as-code

A decision in the wiki, written so a reviewer can check it.

STEP 5

Pre-flight review

Session 2 plans against the wiki before it builds anything.

TAKE HOME

Five steps for Monday

TALK REPO

1. Own identity; MS Learn and a wiki in .mcp.json
2. CLAUDE.md: record every decision, then /build
3. /review: reviewers with their own context and identity
4. /decide the open questions, then sign the baseline
5. /compile-wiki, then @wiki/index.md for the next session

SENTENCE OF THE DAY

Agents make good decisions when you let them record, check and hand them on.

You decide what becomes a rule.

Ticino Data Conference 2026

Thank you to our sponsors

Hosted by

acer

Gold

codeway

Lodestar

IT'S TIME TO INNOVATE

VARGROUP

Silver

Bi Factory

DATA KNOWLEDGE ADVISOR

CLOUDMETRIX

EY

K4ANALYTICS®

PLAN ANYWHERE

OKVIZ®

Thank you

Questions now, or later on LinkedIn

Matthias Falland

linkedin.com/in/matthias-falland

LINKEDIN

TALK REPO