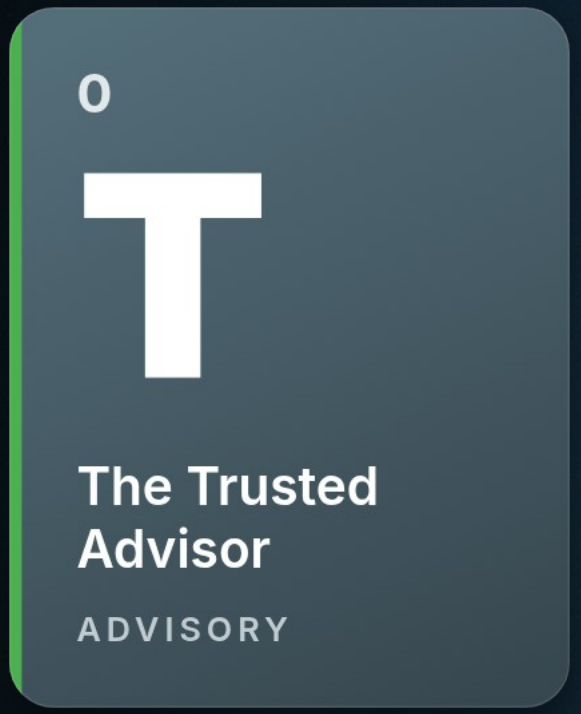

TICINO DATA CONFERENCE 2026 · BIOGGIO

Boss Your Agents: How AI Secures What AI Teams Build

Matthias Falland

MICROSOFT DATA PLATFORM MVP · THE TRUSTED ADVISOR

ABOUT ME

Matthias Falland

The Trusted Advisor · Microsoft Data Platform MVP since 2018

- **Microsoft MVP** Data Platform, since 2018
- **Fabric Friday** weekly on YouTube, one Fabric element per episode
- **Fabric Architecture Podcast** weekly, with an AI co-host
- **Fabric Periodic Table** fabricperiodictable.com, every Fabric item
- **Meetups** Zurich, Hamburg, Basel: over 5,000 members

LinkedIn: matthias-falland
Repo of this talk →

What an agent session is

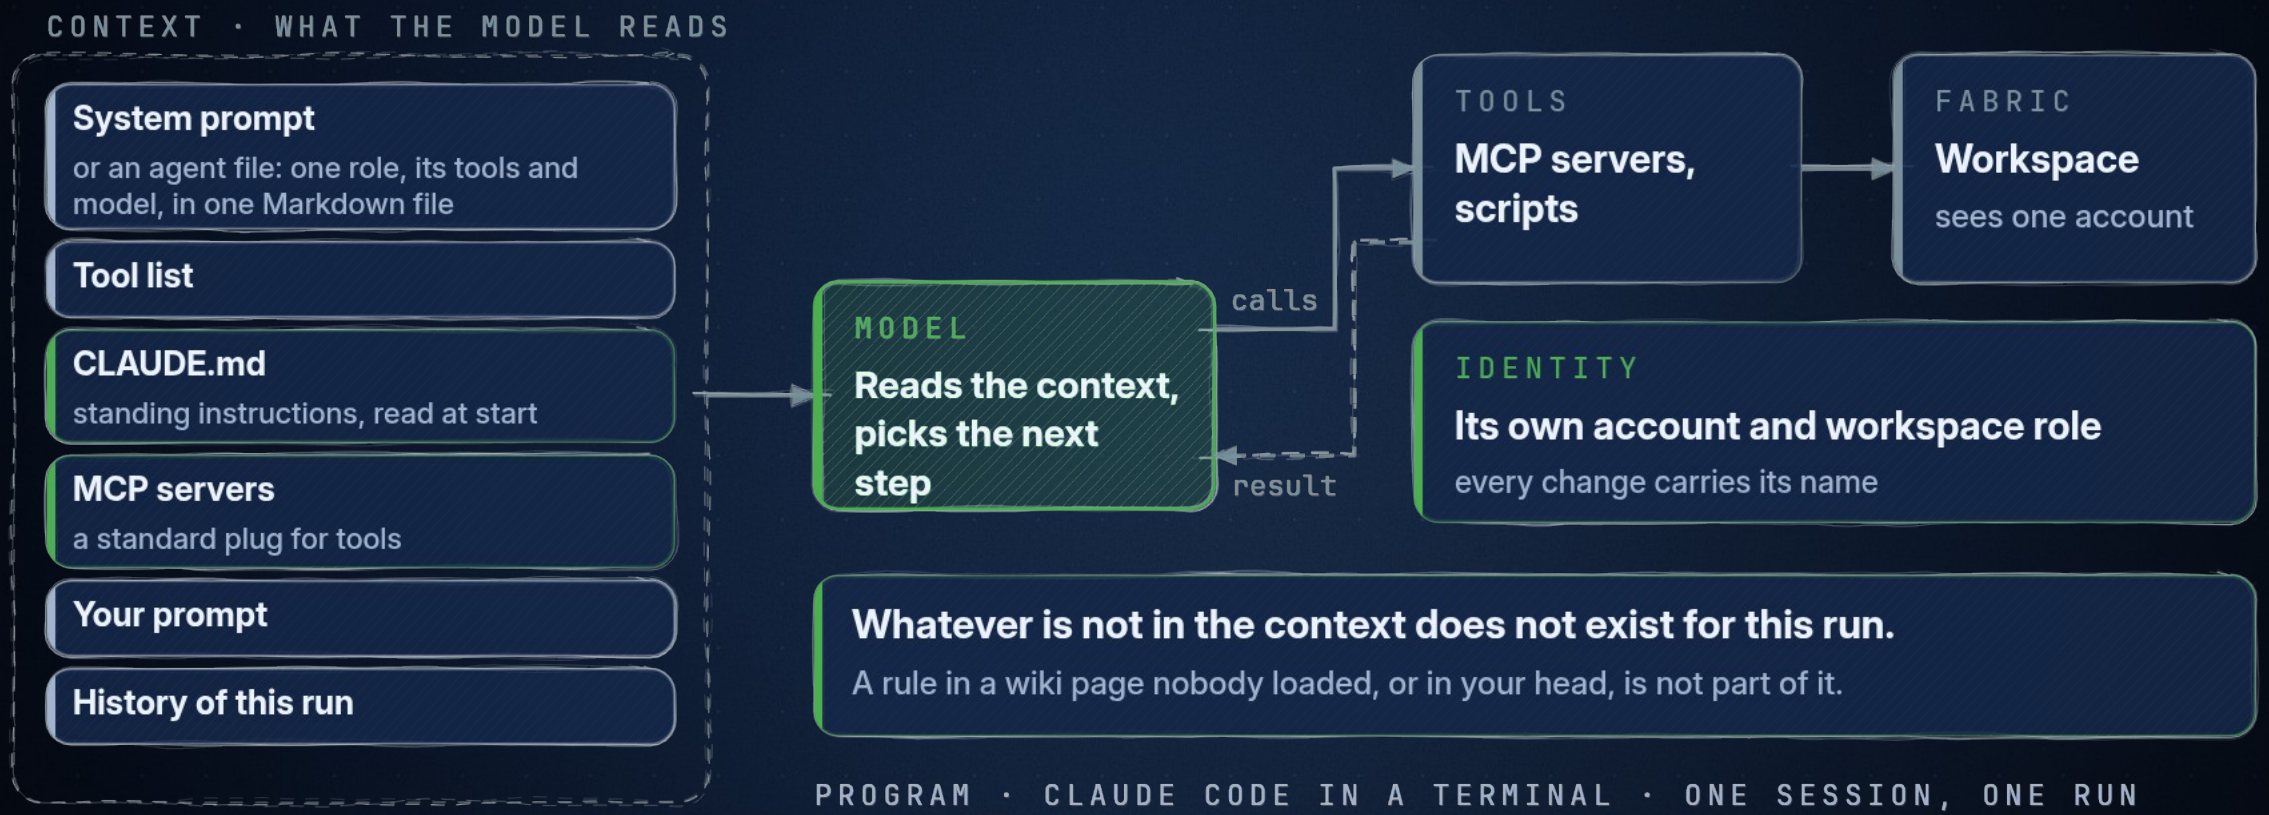

What a project command is

WHAT

A Markdown file in .claude/commands/. The file name becomes the command: build.md is /build.

WHAT FOR

One recurring step as a fixed prompt: the same in every run, shared with the team through the repo.

HOW

Frontmatter: a description, and disable-model-invocation: true so only you start it. $ARGUMENTS takes what you type after the command.

Same thing as a skill:
.claude/skills/build/SKILL.md

SESSION-1/.CLAUDE/COMMANDS/BUILD.MD

```
---
description: Deliver the request in task/.
disable-model-invocation: true
---
Read `task/sales-request.md` and deliver it.
```

THE COMMANDS OF THIS DEMO

/build	Deliver the request in task/.
/review	… Three reviewers read the deployed state …
/decide	Record a decision of the project owner …
/remediate	Apply what a human decided …
/compile-wiki	Compile the project wiki …
/governed-build	… the gate decides, a human approves.

00

Why rules go missing

context, not intelligence

- **The model follows what it reads: your rules are in a wiki and in people's heads**
- **Without them: missing row-level security, inconsistent names, undocumented decisions**
- **A setup and five steps put the project's knowledge into the context, and check it**

THE ARC

Zero to hero: a setup and five steps

THE MAIN CHARACTER IS THE PROJECT'S KNOWLEDGE

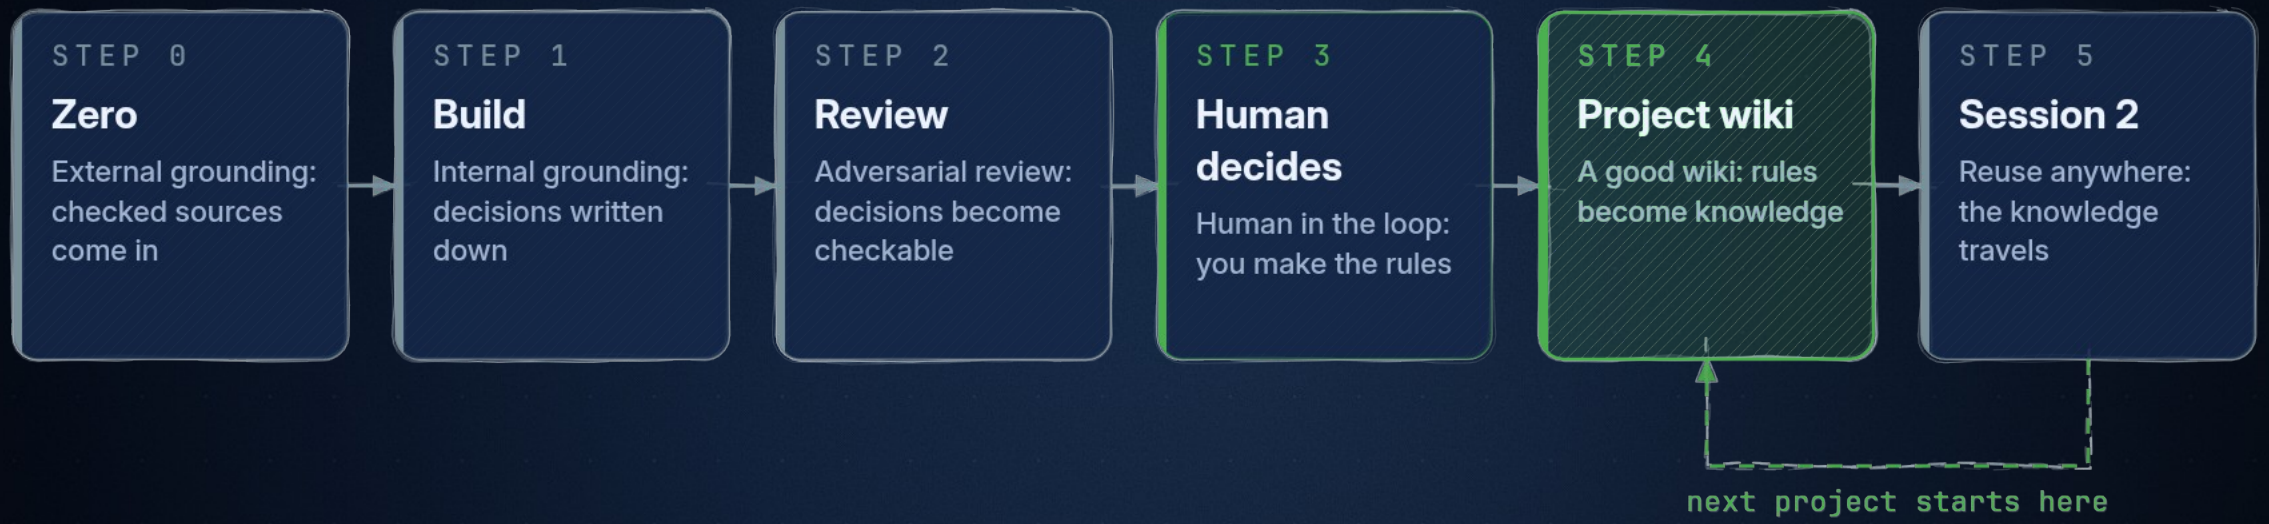

THEN MONDAY MORNING · DRIFT DETECTION: THE KNOWLEDGE DEFENDS ITSELF

NOTHING SEEDED · EVERY STEP IS A FILE OR A COMMAND YOU CAN REDO

Zero

37

WS

Workspace

GOVERN

Where does a session's knowledge come from, and what may it touch?

STEP 0 · PRINCIPLE

External grounding

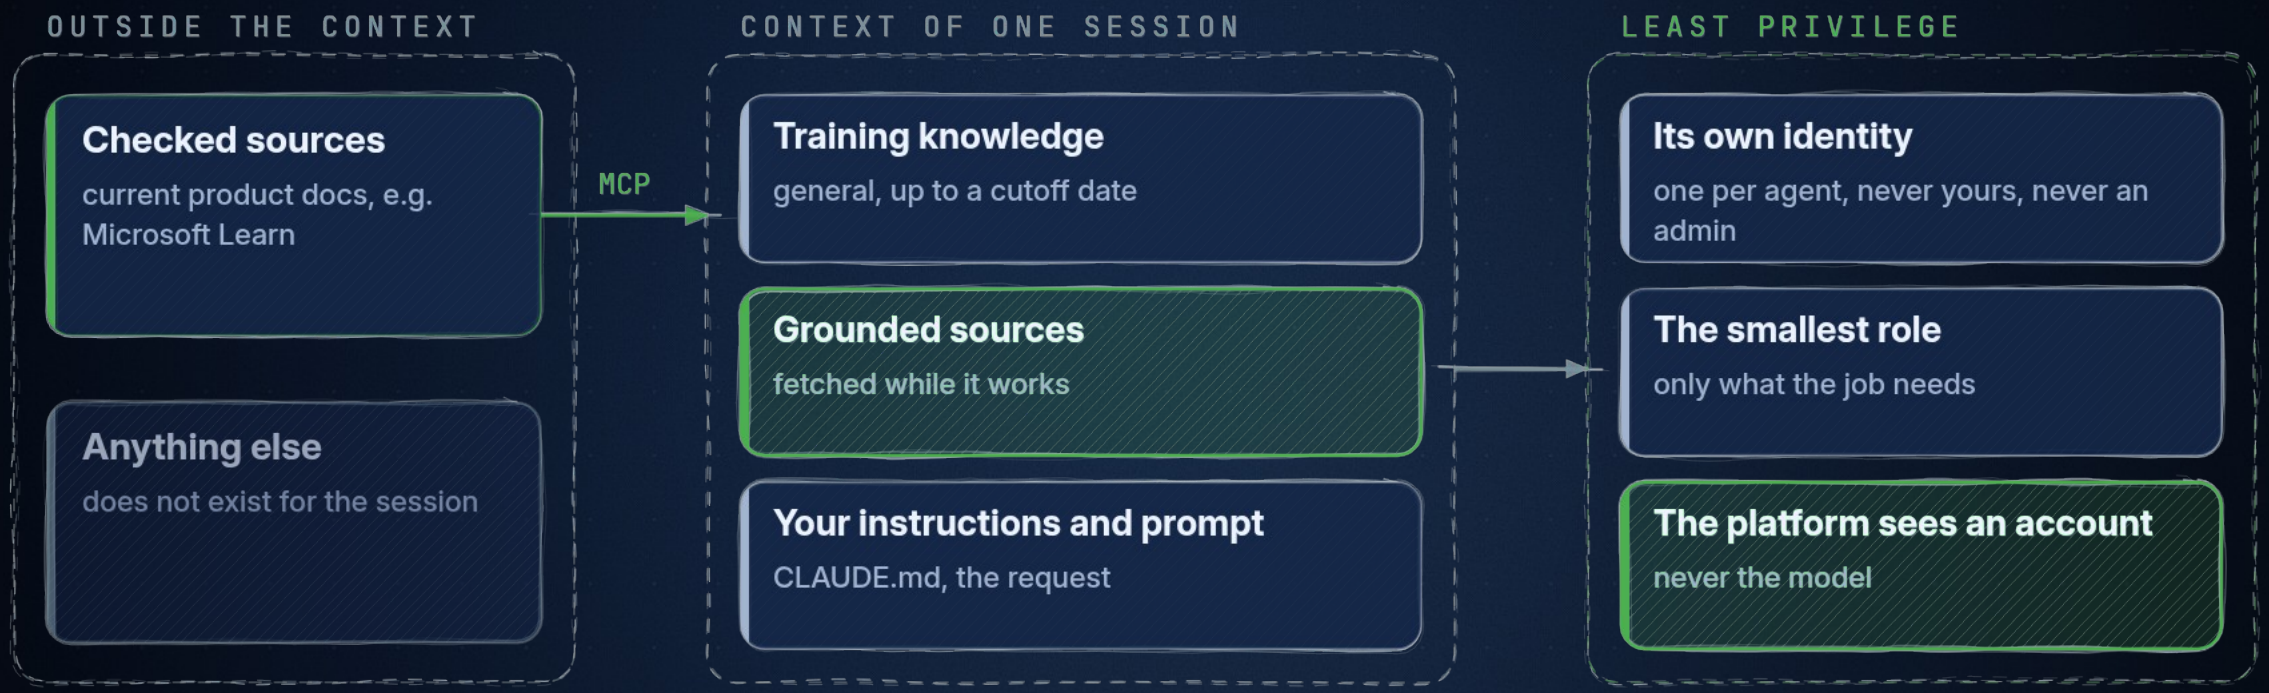

WHATEVER IS NOT IN THE CONTEXT DOES NOT EXIST FOR THE SESSION

FACT CARD · STEP 0

Zero

WHAT

Before the first prompt, each agent gets its own identity, the smallest role, Microsoft Learn and a wiki (a folder in the repo is enough).

GOOD WHEN

You start a project, or hand an existing one to an agent for the first time.

IN YOUR PROJECT

```
Identity per agent
  never yours, never an admin
Microsoft Learn MCP, a wiki
CLAUDE.md: log every decision:
  what, why, alternatives, source
Deny the builder access
  to the reviewers' files
.claude/commands/: a prompt
  file per step
```

TAKEAWAY

A decision log needs an instruction and a place to write.

1 Build

How does a decision outlive the session that made it?

STEP 1 · PRINCIPLE

Internal grounding

DEMO · HOW IT RUNS

One real run, recorded in one go

- Microsoft Fabric, an F2 capacity, two empty workspaces
- Claude Code; a small toolkit calls the Fabric REST API with each agent's own account
- Waiting time is cut; the right column says what to look for
- From step 3 on, live from the state this run left

Step 1 · Session 1 builds and records its decisions

Direct Lake on OneLake (not on SQL endpoint)

Decision: The semantic model `Sales` uses Direct Lake on OneLake: shared expression `AzureStorage.DataLake("https://onelake.dfs.fabric.microsoft.com/<ws id>/<lakehouse id>")`, partitions `mode: directLake`.

Why: Microsoft recommends Direct Lake on OneLake for new semantic models; it never falls back to DirectQuery, and semantic-model RLS is supported.

Alternatives: Direct Lake on SQL analytics endpoint (can silently fall back to DirectQuery); Import (requires data refresh copies, not what was asked).

Source: **https://learn.microsoft.com/fabric/fundamentals/direct-lake-how-it-works#directquery-fallback** https://learn.microsoft.com/fabric/fundamentals/direct-lake-how-it-works#directquery-fallback , **https://learn.microsoft.com/fabric/fundamentals/direct-lake-develop#power-query-connectors** https://learn.microsoft.com/fabric/fundamentals/direct-lake-develop#power-query-connectors

(END)

④ Takeaway

Read the pages yourself. A good page says what was decided, why, what else was possible, and where it is documented.

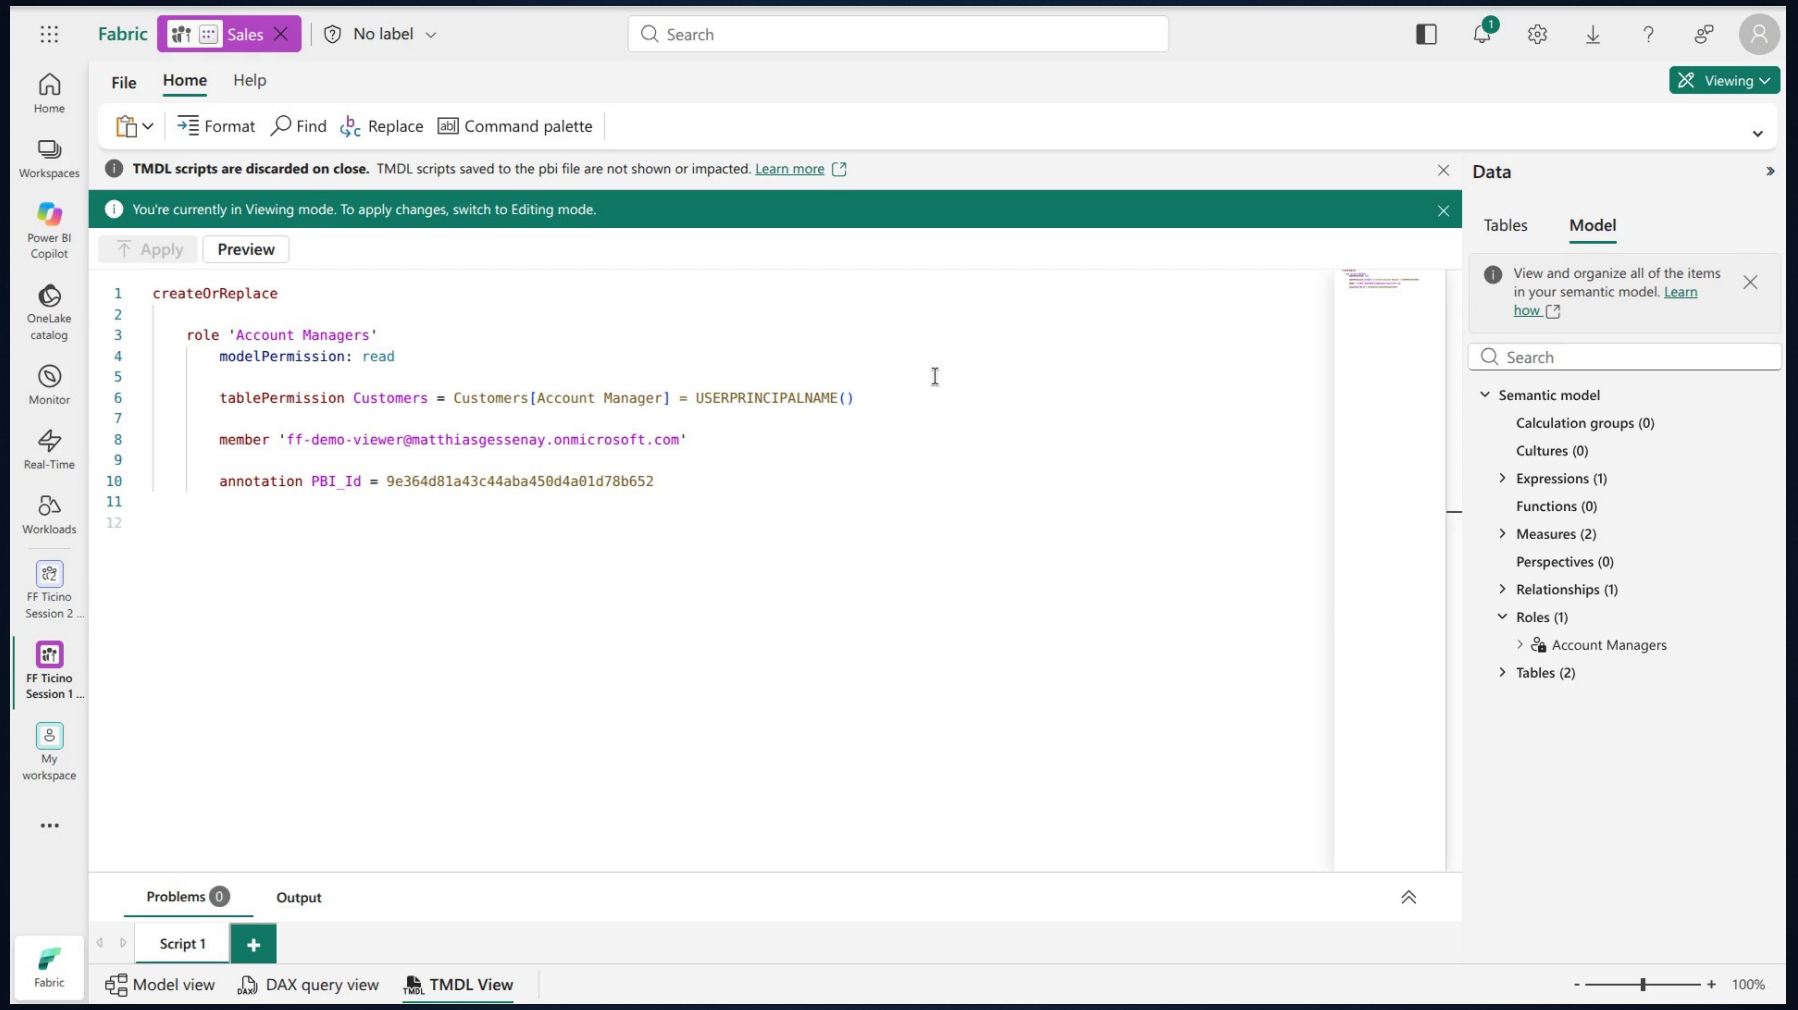
Fabric
Sales
No label
Search
File
Home
Help
Viewing
Format
Find
Replace
Command palette
TMDL scripts are discarded on close. TMDL scripts saved to the pbi file are not shown or impacted. Learn more
You're currently in Viewing mode. To apply changes, switch to Editing mode.
Apply
Preview
createOrReplace
role 'Account Managers'
modelPermission: read
tablePermission Customers = Customers[Account Manager] = USERPRINCIPALNAME()
member 'ff-demo-viewer@matthiasgessenay.onmicrosoft.com'
annotation PBI_Id = 9e364d81a43c44aba450d4a01d78b652
Data
Tables
Model
View and organize all of the items in your semantic model. Learn how
Search
Semantic model
Calculation groups (0)
Cultures (0)
Expressions (1)
Functions (0)
Measures (2)
Perspectives (0)
Relationships (1)
Roles (1)
Account Managers
Tables (2)
Problems 0
Output
Script 1
Model view
DAX query view
TMDL View
100%
Home
Workspaces
Power BI Copilot
OneLake catalog
Monitor
Real-Time
Workloads
FF Ticino Session 2 ...
FF Ticino Session 1 ...
My workspace
Fabric

FACT CARD · STEP 1

Build with a decision log

WHAT

The builder works on the request and writes one wiki page per decision while it builds.

GOOD WHEN

An agent makes a choice that someone will ask about later.

IN YOUR PROJECT

The request as written
  no hidden rules
One wiki page per decision
  what · why ·
  alternatives · source
Written when decided,
  not at the end

THE PROMPT BEHIND /build

```
.claude/commands/build.md
Read `task/sales-request.md` and deliver it.
```

TAKEAWAY **A good page says what, why, what else was possible, and the source.**

2 Review

Who checks the builder, and against what?

STEP 2 · PRINCIPLE

Adversarial review

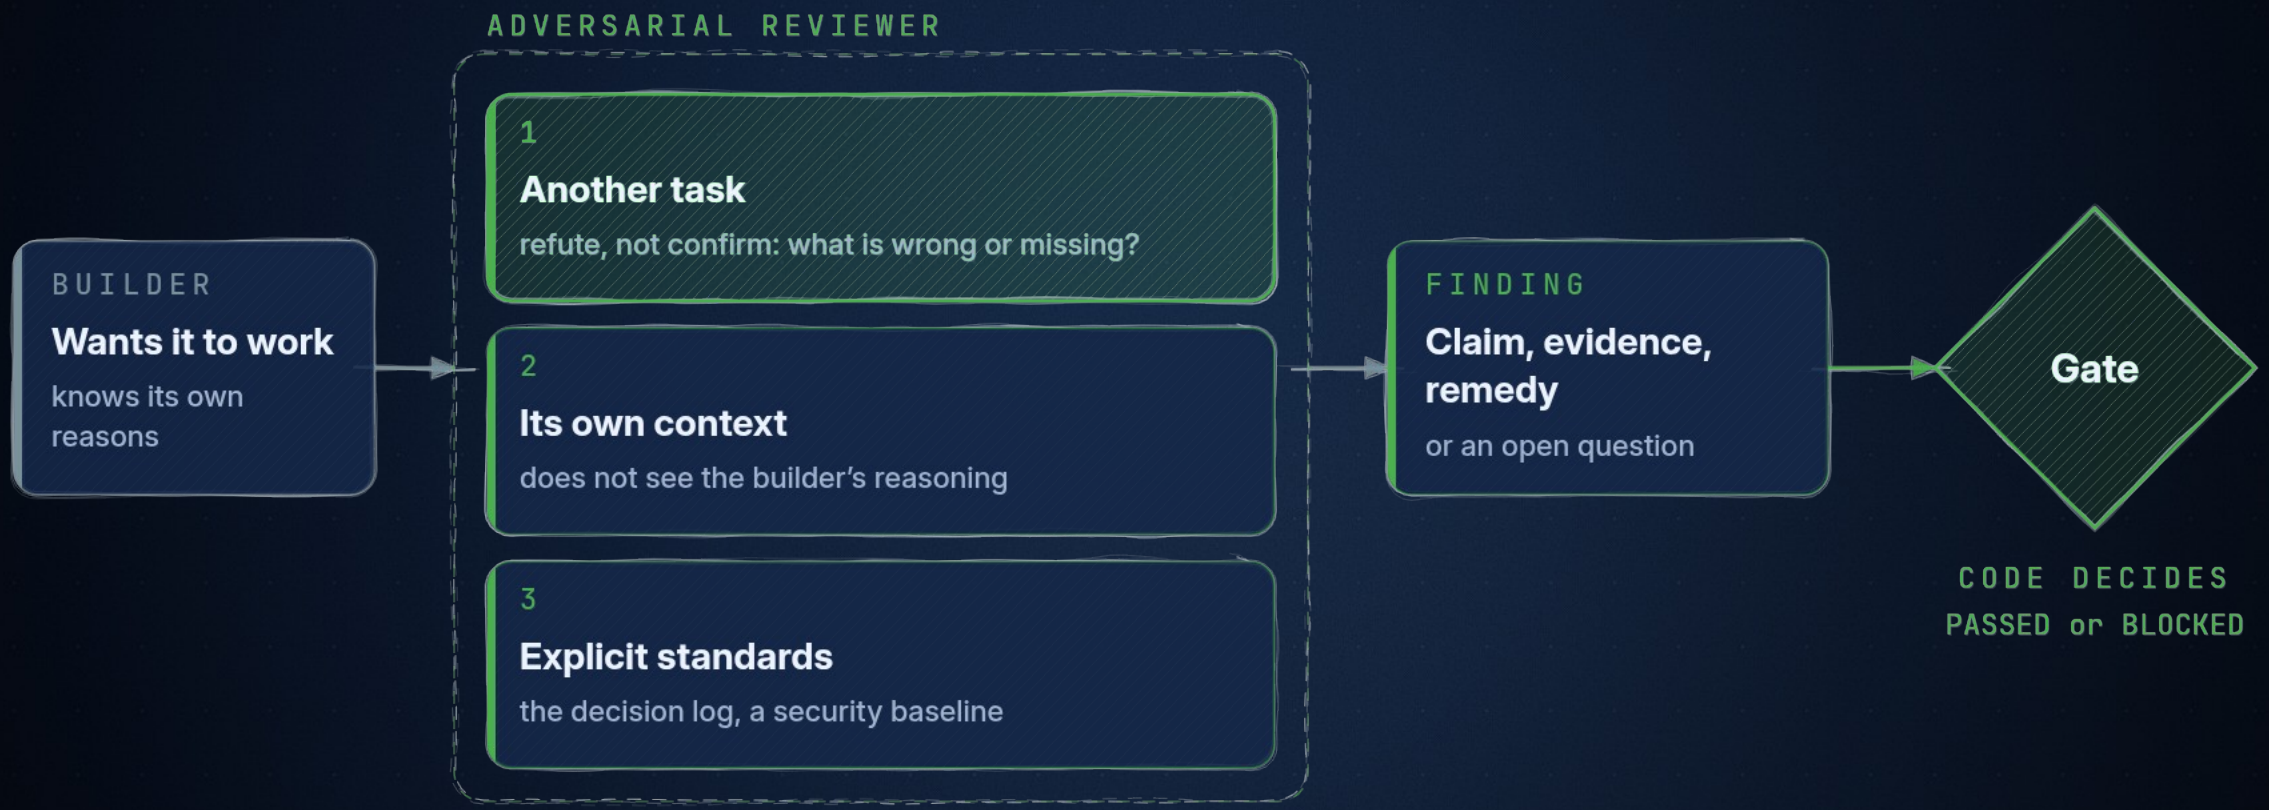

STEP 2 · EVIDENCE

A finding is a claim, evidence and a remedy

CLAIM

What the reviewer says

one sentence, checkable

EVIDENCE

What was read, by whom

item, definition, query · which identity

REMEDIATION

What would close it

a change, or a decision

SEC-020 · BLOCKING · FROM THE REVIEW

CLAIM

RLS bypass: a role member reads 60 of 60 rows through the SQL endpoint

EVIDENCE

A test identity in the RLS role read the customers table through the SQL endpoint

60 rows, 40 of them belong to other account managers

REMEDIATION

Apply the same rule wherever the table can be read

a security policy in the SQL endpoint, or OneLake security

MODEL RLS APPLIES ONLY INSIDE THE MODEL

Step 2 · Three reviewers read the real state

```
$ sed -n '/^## Open/,/^## Info/p' reports/latest-gate.md | glow -w 84 -

  ## Open questions (a human decides)

  ### Q-1 · open · how items are named

  What naming convention should items follow (why sales_lh and Sales, no prefix/suffix
  or environment tag)?

  • Evidence: F1: items built: Lakehouse «sales_lh»; SemanticModel «Sales». Names
  appear only on pages about other choices. No decision page covers it.
  • Remediation: a human decides (one sentence) and the decision goes into the log
  • Reviewer: decisions

  ## Info

$
```

④ Takeaway

Test what a rule does, not only that it exists. A choice nobody made, like item names, becomes a question for a person.

FACT CARD · STEP 2

Review with a fresh context

WHAT

Reviewers with their own identity and a fresh context read the deployed state and the decision log, and test what a role member really sees.

GOOD WHEN

Before anyone builds on a result, and after every change.

IN YOUR PROJECT

One agent file per reviewer
  one question each
Its own identity
  reads the real state
A test identity in the role
  reads through every engine
A finding:
  claim · evidence · remedy
No decision? an open question

THE PROMPT BEHIND /review

.claude/commands/review.md

```
You are the orchestrator. You do not review
and you do not fix; you route, and you stop
at the gate.

1. Run `./py gate/new_run.py`.
2. Use the **security-reviewer**,
   **decision-reviewer** and
   **drift-reviewer** agents **in parallel**.
3. Run `./py gate/gate.py`. Show the verdict
   and every finding, … Do not explain
   findings away: a finding stays as the
   reviewer reported it.
```

TAKEAWAY

Test the effect: model RLS applies only inside the model.

DEMO · LIVE FROM HERE

The same run, continued live

- Workspace one is as the review found it, nothing rebuilt
- I decide; the agent fixes; the review runs again; I sign
- An agent writes the wiki; session two starts with it
- If the network fails, the recording of the same run takes over

3

Human decides

Who decides what becomes a rule?

STEP 3 · PRINCIPLE

Human in the loop

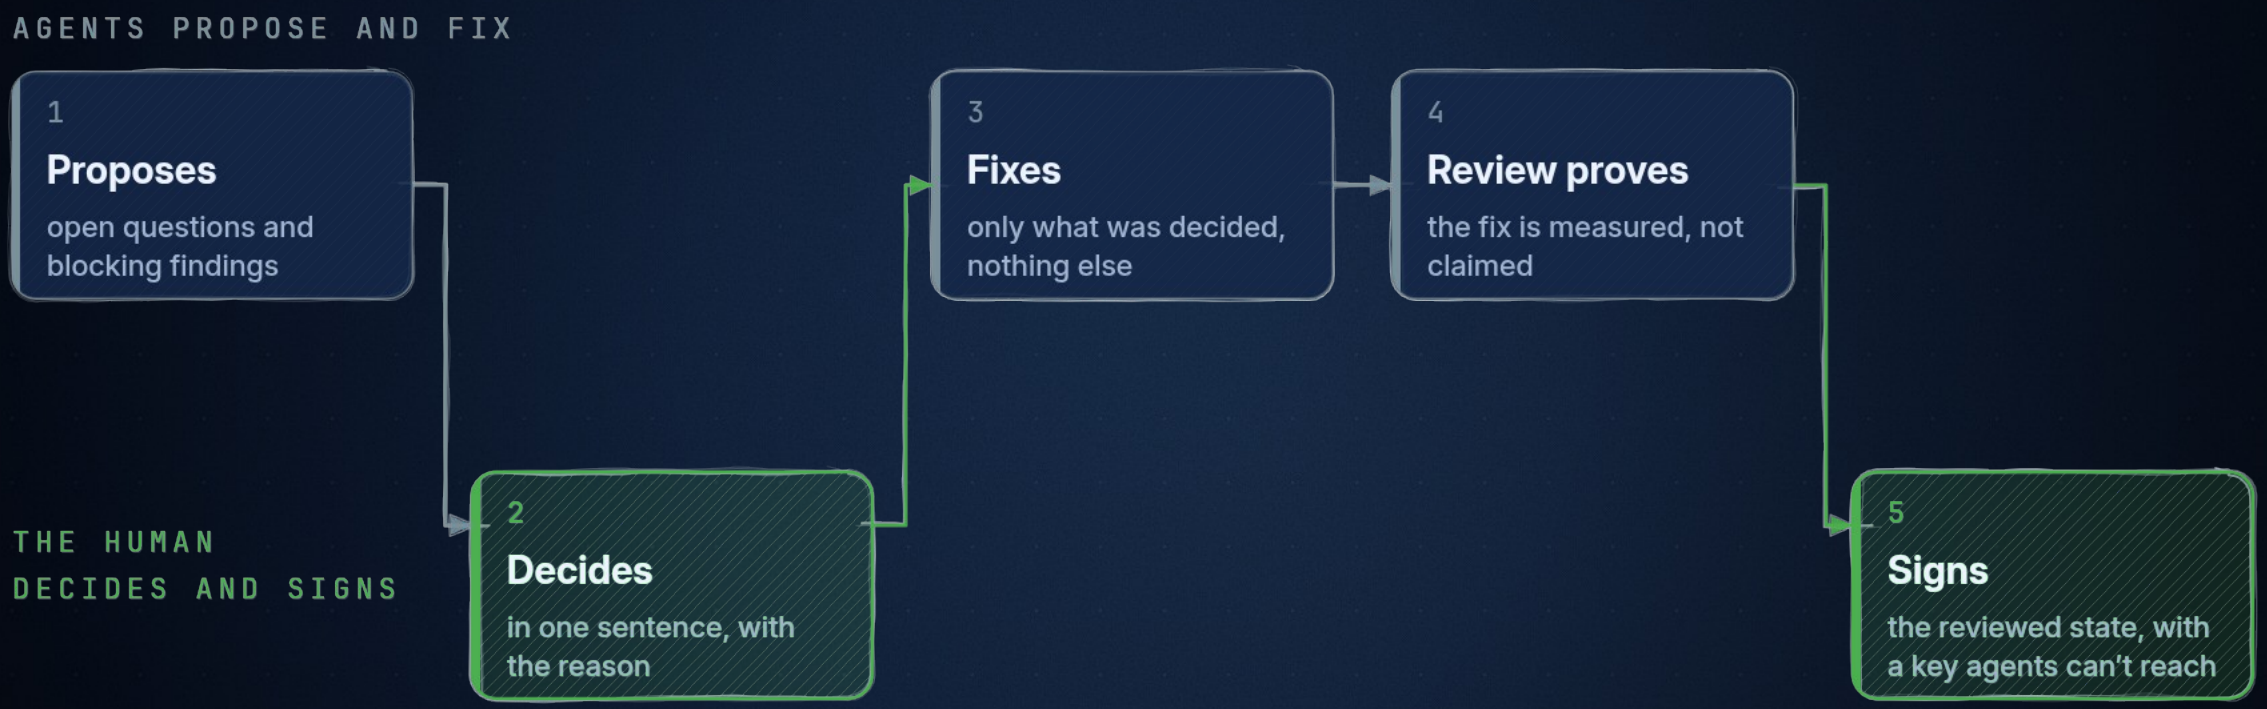

Step 3 · The person decides, the agent remediates, the review proves it

```
$ grep -o 'reads [0-9]* of [0-9]* rows[^|]*' reports/latest-gate.md | fold -s -w 84
reads 20 of 60 rows of [dbo].[customers] in lakehouse sales_lh through its SQL
analytics endpoint (u62solgdfuke...), none of another user, its own rows: 20
$ ./py gate/snapshot.py sign
snapshot signed by Matthias Falland: 2 items (Lakehouse/sales_lh, SemanticModel/Sales),
 8 definition parts -> /home/demo/.local/state/zero-to-hero/·····················
·············/snapshot.json, reports/baseline.json
$
```

④ **Takeaway**

A person decides, an agent remediates, the review measures. The person signs the state in the shell, not the agent.

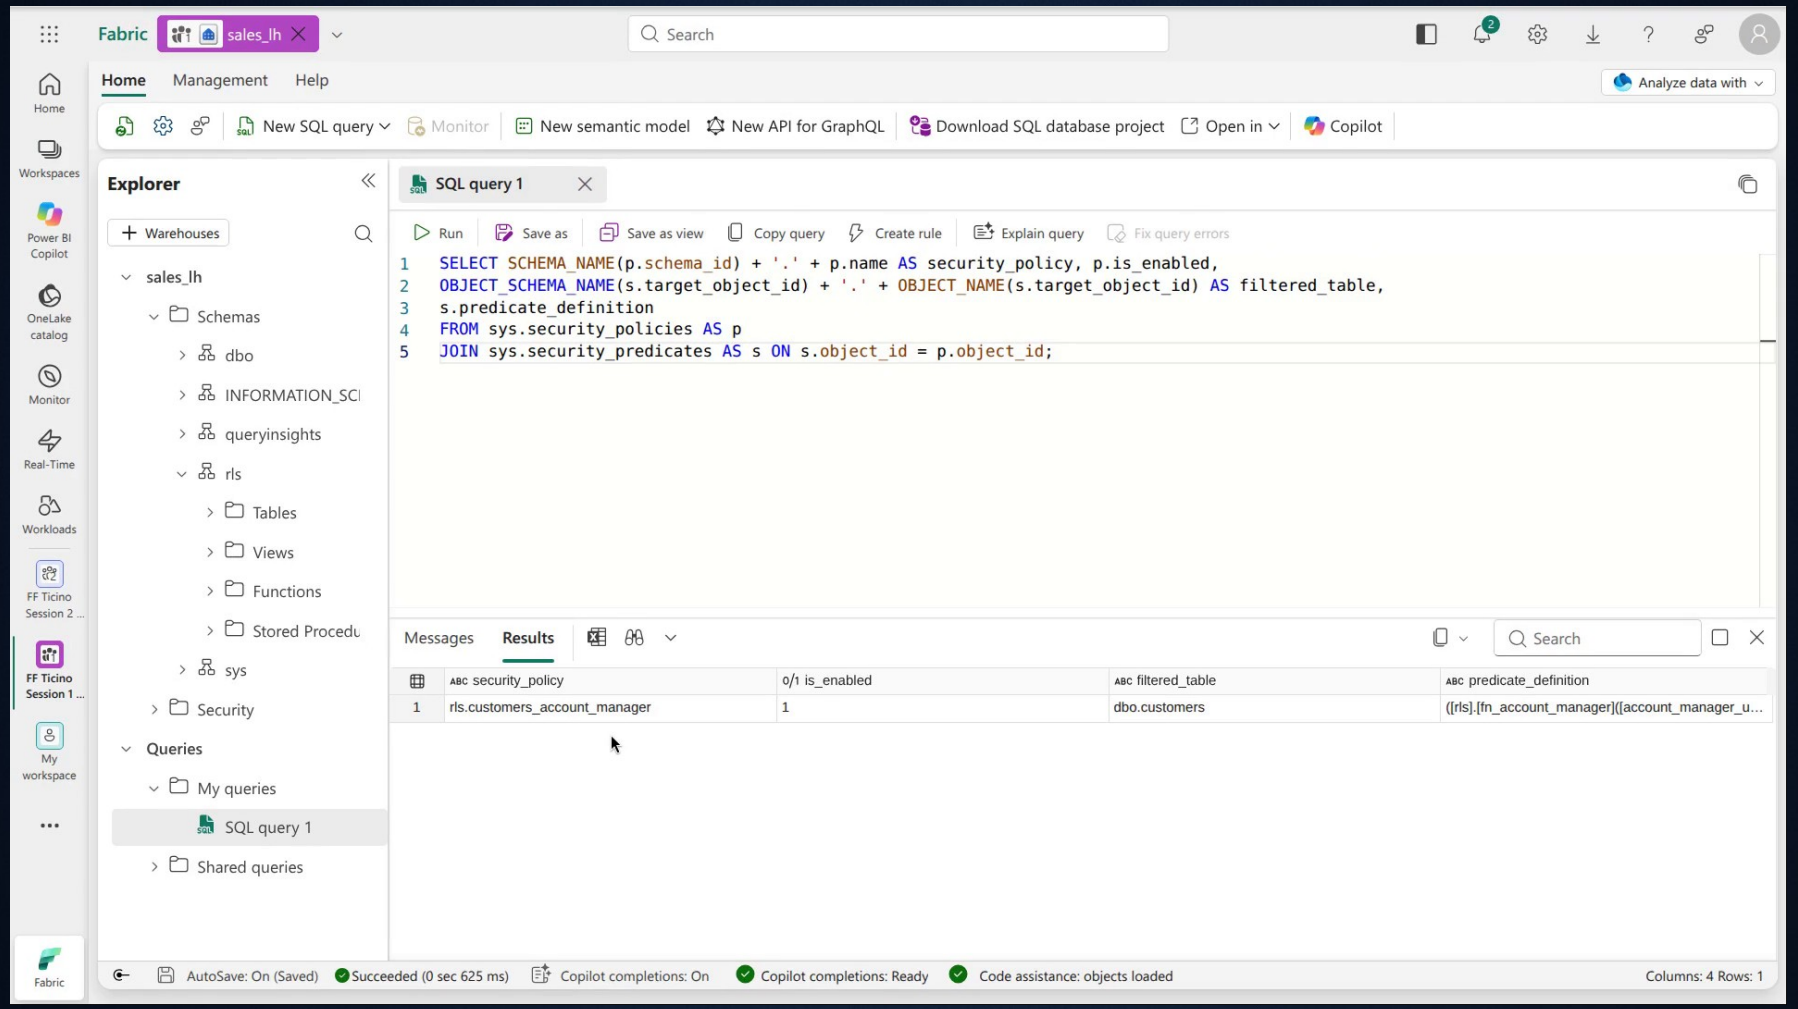
Fabric
sales_lh
Search
Home
Management
Help
Analyze data with
New SQL query
Monitor
New semantic model
New API for GraphQL
Download SQL database project
Open in
Copilot
Explorer
Warehouses
sales_lh
Schemas
dbo
INFORMATION_SCI
queryinsights
rls
Tables
Views
Functions
Stored Procedu
sys
Security
Queries
My queries
SQL query 1
Shared queries
SQL query 1
Run
Save as
Save as view
Copy query
Create rule
Explain query
Fix query errors
1 SELECT SCHEMA_NAME(p.schema_id) + '.' + p.name AS security_policy, p.is_enabled,
2 OBJECT_SCHEMA_NAME(s.target_object_id) + '.' + OBJECT_NAME(s.target_object_id) AS filtered_table,
3 s.predicate_definition
4 FROM sys.security_policies AS p
5 JOIN sys.security_predicates AS s ON s.object_id = p.object_id;
Messages
Results
Search
security_policy
is_enabled
filtered_table
predicate_definition
1
rls.customers_account_manager
1
dbo.customers
([rls].[fn_account_manager]([account_manager_u…
AutoSave: On (Saved)
Succeeded (0 sec 625 ms)
Copilot completions: On
Copilot completions: Ready
Code assistance: objects loaded
Columns: 4 Rows: 1

FACT CARD · STEP 3

The human decides

WHAT

You decide each open question and blocking finding in one sentence; the agent fixes only that, the review runs again, you sign.

GOOD WHEN

A finding blocks, or a choice should become a rule.

IN YOUR PROJECT

```
One sentence per question
  or blocking finding
The agent fixes what
  you decided, nothing else
The review runs again:
  the fix is measured
Then sign the baseline
  key the agents can't reach
```

TWO PROMPTS: /decide, /remediate

.claude/commands/decide.md

```
The project owner, Matthias Falland,
decided: "$ARGUMENTS"
… 3. Record it with `wiki_add` as one page …
   - **Why**: his reason if he gave one; …
     Never invent a reason. …
```

.claude/commands/remediate.md

```
Apply the human's decisions to the
blocking findings of the latest review,
and nothing else.
… Leave alone every finding without a
human decision, every open question and
everything else in the workspace. …
```

TAKEAWAY

You decide, the agent fixes, the review proves.

4

Project wiki

What belongs in a project wiki, and why each page type?

STEP 4 · PRINCIPLE

What a good wiki holds

WIKI/ · A FOLDER IN THE REPOSITORY

index.md	**The entry point** — one line per page, read first
ADR-*	**Architecture Decision Record** — context, decision, alternatives, consequences
SEC-*	**Security finding or rule** — with its measurement
DRF-*	**Drift baseline** — the signed target state to compare with
rules/	**Checkable rules** — distilled from ADRs, a script checks them

Few page types
each with a clear purpose

Reviewed like code
a pull request before it merges

Read by the next session
and by people

Step 4 · From log, report and decisions to the wiki

ADR-005 Item naming with type prefixes

Status

Accepted. Decided by the human: Matthias Falland, project owner, 2026-10-28 (tag `decided-by-human`). Applies from the next delivery.

Context

Open question Q-1 of the gate review of 2026-10-28 07:19 asked: "What naming convention should items follow (why sales_lh and Sales, no prefix/suffix or environment tag)?" The items had been built as Lakehouse «sales_lh» (type as suffix) and SemanticModel «Sales» (no type tag, capitalised). No convention had been decided.

Decision

"Items use type prefixes: lh_, sm_, rpt_, nb_, lower case, underscores, so anyone can tell an item's type from its name."

Why

So anyone can tell an item's type from its name. Consistent, predictable names help users and administrators.

Alternatives considered

Keep the names as deployed (`sales_lh`, `Sales`): type as suffix or missing, mixed case, no environment tag.

:

④ Takeaway

A rule belongs in the wiki when a reviewer can check it. Everything else stays a recorded decision.

FACT CARD · STEP 4

The project wiki

WHAT

Log, reviews and your decisions become one folder. Rules are data a script checks; the gate verdict is code.

GOOD WHEN

A project ends, or the next team or agent starts.

IN YOUR PROJECT

```
wiki/index.md, ADR-*, SEC-*, DRF-*
wiki/rules/ADR-005.yaml
  id: ADR-005
  status: accepted
  decided_by: the human
  check:
    item_names:
      SemanticModel: '^sm_…'
```

THE PROMPT BEHIND /compile-wiki

```
.claude/commands/compile-wiki.md
Compile the project wiki of this delivery.
Your inputs, and nothing else:
- the decision log: every page in the wiki …
- the review: every gate report …
- the signed baseline …
1. **ADR pages**, one per accepted decision …
2. **SEC pages**, one per security finding …
3. **DRF page**, one per drift finding …
4. **Rules** for the next review: …
5. Run `./py ../review/gate/check_rules.py
   --lint rules` and fix the files until it
   reports 0 errors.
…
```

TAKEAWAY

A rule belongs in the wiki when a reviewer can check it.

5
Session 2

How does the next session, or a colleague, start with what we learned?

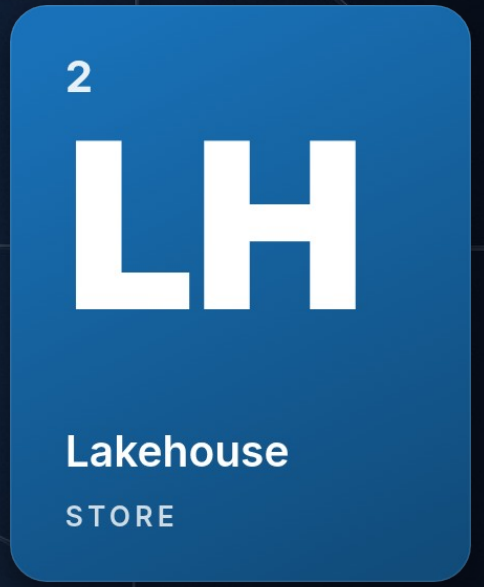

STEP 5 · PRINCIPLE

Reuse anywhere

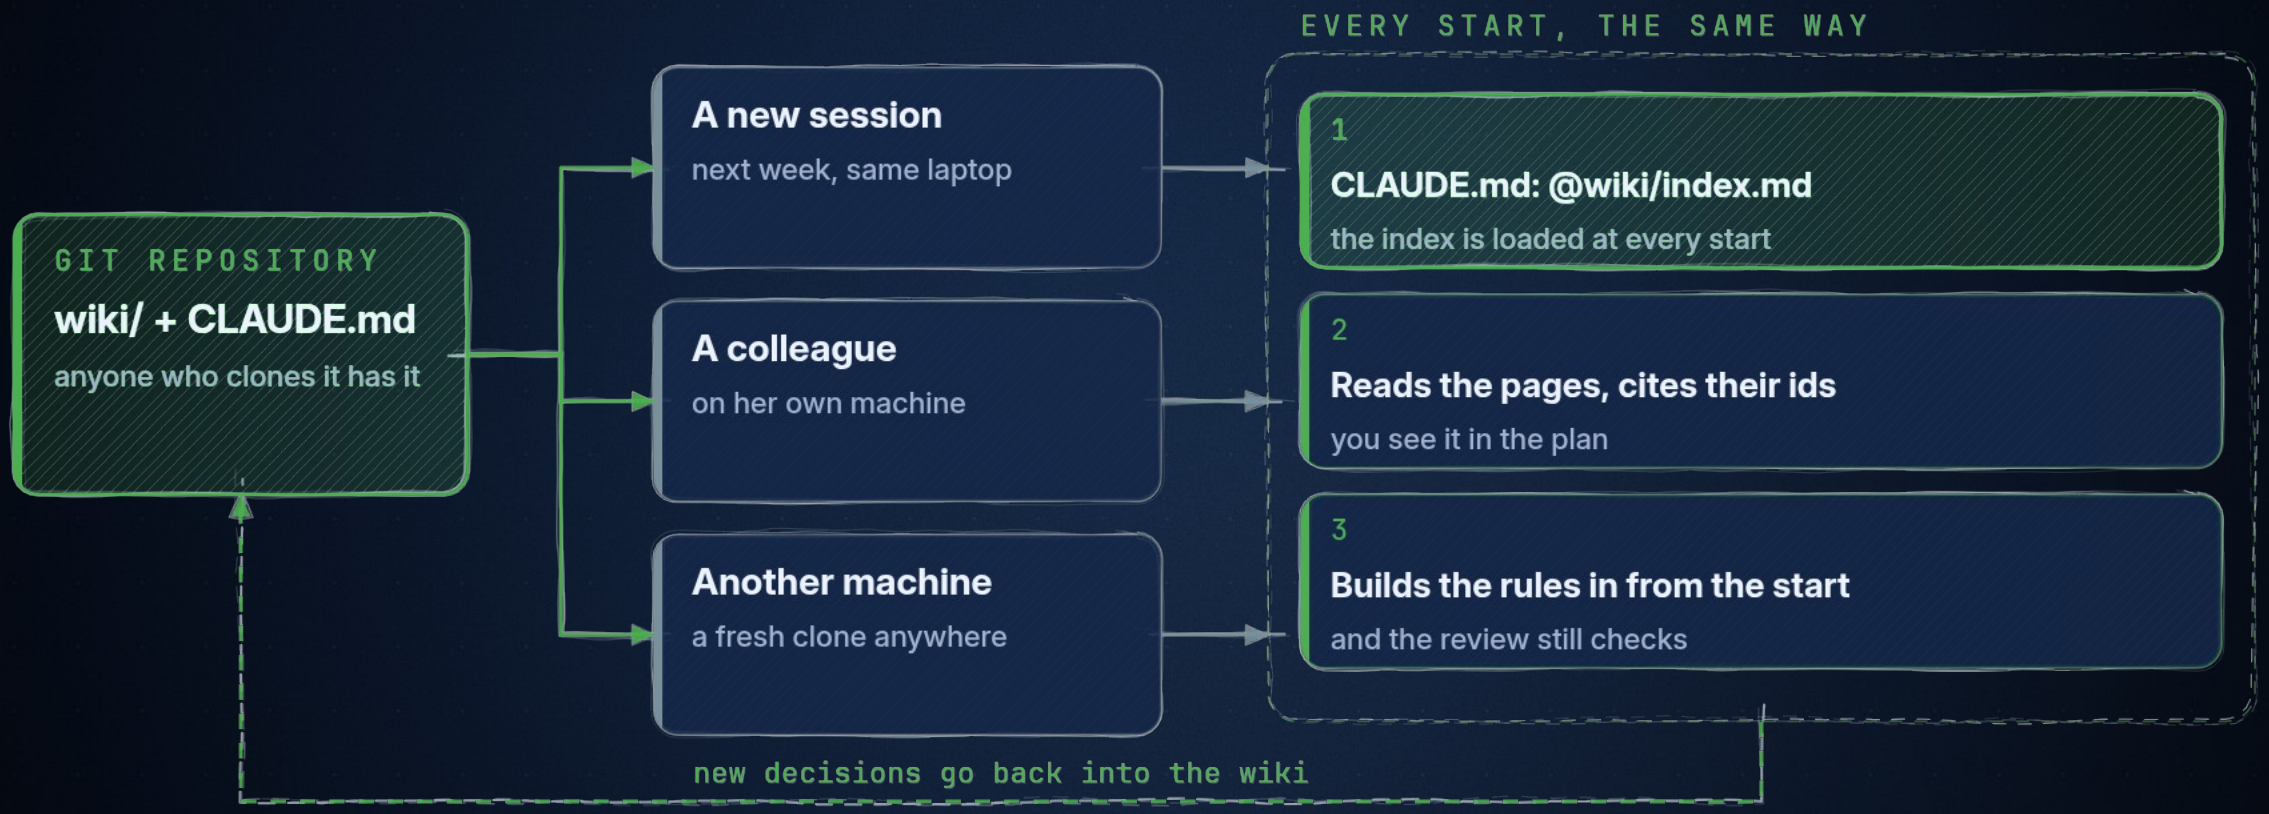

Step 5 - Session 2 plans with the wiki, the review confirms

Plan: Sales request (Giulia, Sales Operations) - session 2

Workspace: FF Ticino Session 2 DEV (a12f95a1-...), empty at start. Identity: ff-demo-contributor (builder).
Request: `task/sales-request.md`. Data: `data/customers.csv` (60 rows, 3 account managers x 20 customers),
`data/orders.csv` (600 rows; all 60 customer ids exist in customers).

Wiki read before planning: ADR-001..ADR-006, the 8 decision pages, SEC-010, SEC-020,
DRF-000, and
`wiki/rules/` (ADR-001, -002, -004, -005, -006 yaml).

Items

Item	Type	Name	Follows
Lakehouse	Lakehouse	`lh_sales` (no schemas)	ADR-005 (type prefixes, a…
Tables	Delta	`customers`, `orders` unde…	ADR-001
Semantic mo…	SemanticMod…	`sm_sales`	ADR-005, ADR-002, ADR-003…
SQL endpoin…	T-SQL objec…	schema `rls`, TVF `rls.fn_…`	ADR-006, SEC-020

Model `sm_sales` (definition in `work/sm_sales/`):

• Direct Lake on OneLake: shared expression `AzureStorage.DataLake("https://onelake.dfs.fabric.microsoft.com/<ws>/<lh>")`,
partitions `mode: directLake`, no `schemaName` (lakehouse without schemas) (ADR-002).
• Star schema (ADR-003): `Customers` (Customer Name, Region, Account Manager, Credit Limit; key hidden),
:

④ Takeaway

The plan in `work/PLAN.md` cites the ADRs and SEC-020 before anything is built. Hand the knowledge on, and keep the review.

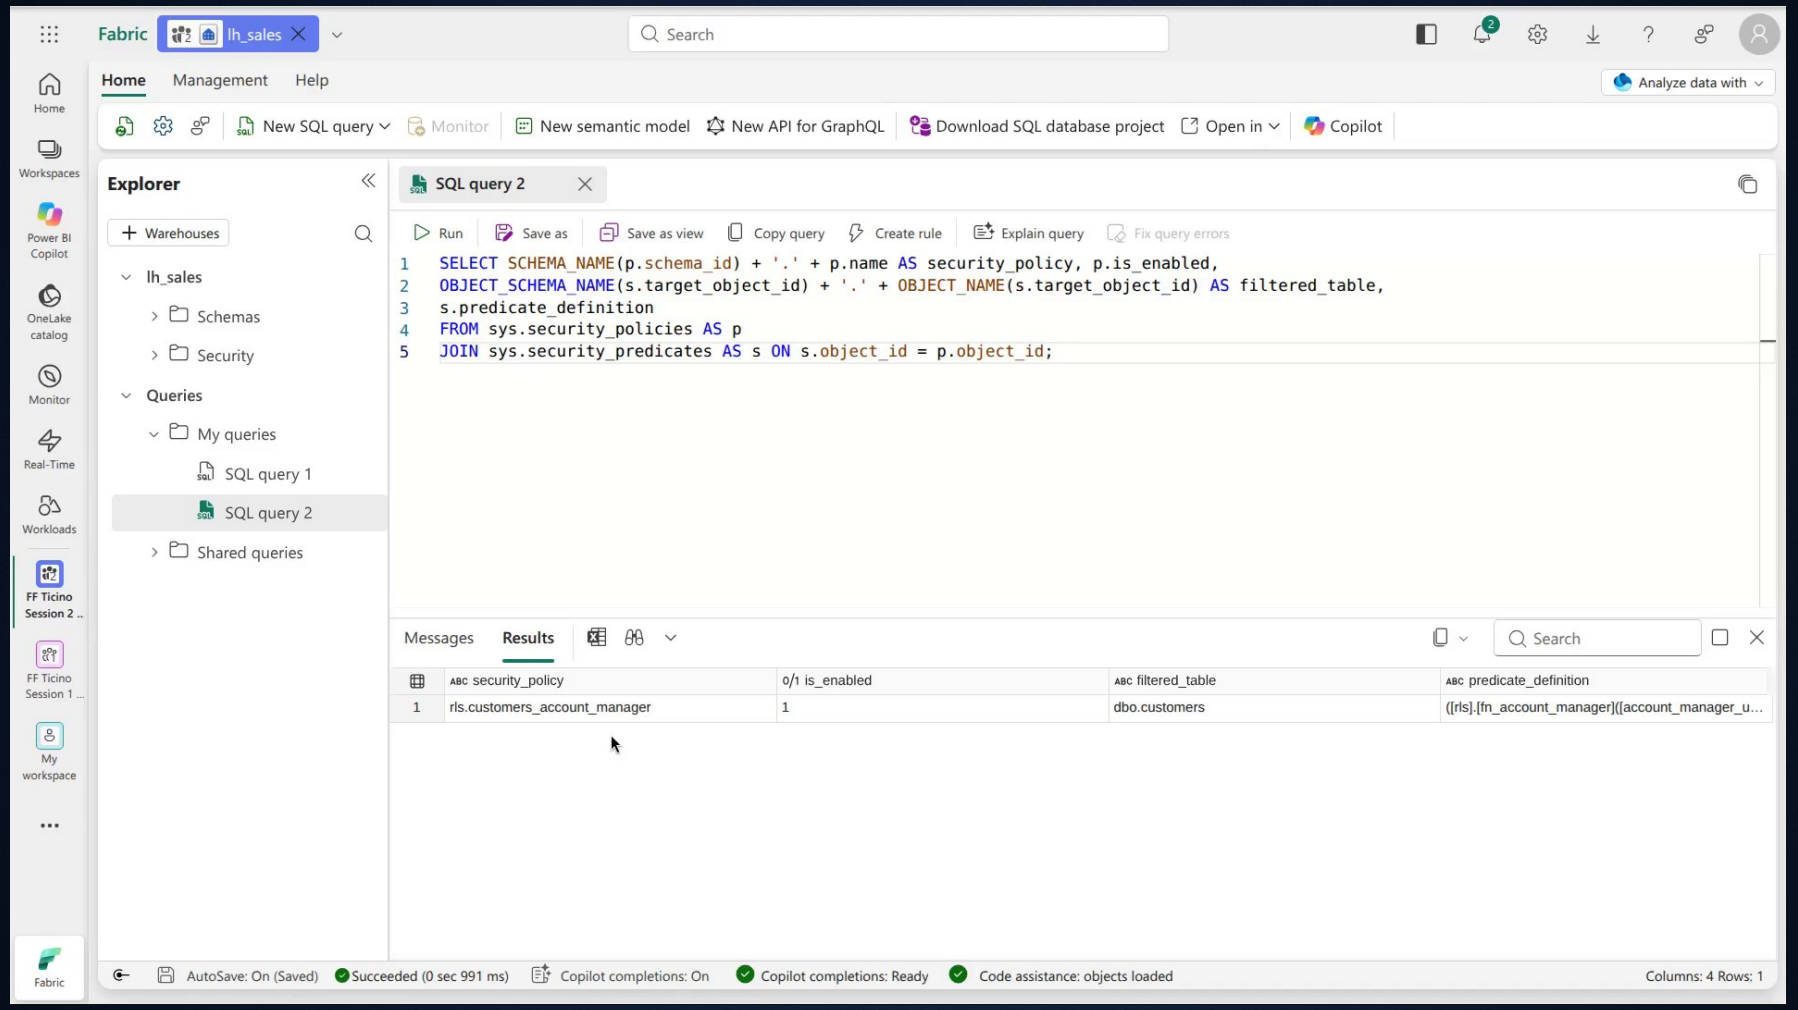
Fabric
lh_sales
Search
Home
Management
Help
Analyze data with
New SQL query
Monitor
New semantic model
New API for GraphQL
Download SQL database project
Open in
Copilot
Explorer
Warehouses
lh_sales
Schemas
Security
Queries
My queries
SQL query 1
SQL query 2
Shared queries
SQL query 2
Run
Save as
Save as view
Copy query
Create rule
Explain query
Fix query errors
SELECT SCHEMA_NAME(p.schema_id) + '.' + p.name AS security_policy, p.is_enabled,
OBJECT_SCHEMA_NAME(s.target_object_id) + '.' + OBJECT_NAME(s.target_object_id) AS filtered_table,
s.predicate_definition
FROM sys.security_policies AS p
JOIN sys.security_predicates AS s ON s.object_id = p.object_id;
Messages
Results
Search
security_policy
is_enabled
filtered_table
predicate_definition
1
rls.customers_account_manager
1
dbo.customers
([rls].[fn_account_manager]([account_manager_u…
AutoSave: On (Saved)
Succeeded (0 sec 991 ms)
Copilot completions: On
Copilot completions: Ready
Code assistance: objects loaded
Columns: 4 Rows: 1

FACT CARD · STEP 5

Session 2 starts with knowledge

WHAT

A new session reads the wiki, plans with its ids and builds what earlier findings taught from the start. Reviewers check the result; a human approves.

GOOD WHEN

Every new project, and every new agent on an old one.

IN YOUR PROJECT

```
One line in CLAUDE.md:
  @wiki/index.md, read the
  pages it lists, cite ids
A plan before the build
  cites the ADRs and SEC pages
Reviewers: security,
  rules, drift
A human approves
```

THE PROMPT BEHIND /governed-build

.claude/commands/governed-build.md

```
You are the orchestrator. You do not build …
2. Use the **builder** agent to deliver the
   request. …
3. Use the **security-reviewer**,
   **rules-reviewer** and **drift-reviewer**
   agents **in parallel**.
4. Run `./py gate/gate.py`. …
5. If BLOCKED: … At most two remediation
   rounds; …
6. If PASSED: ask the human whether to
   approve. Only after an explicit
   "approve", run
   `./py gate/snapshot.py sign`. …
```

TAKEAWAY

Hand the knowledge on, and keep the review.

Monday morning

40

DPL

Deployment Pipeline

GOVERN

Drift detection: someone renames the model by hand. Does the project notice?

Step 6 · Monday morning

● Agent "Rules review of deployed state" finished · 58s

› Message from @security-reviewer (ctrl+o to expand)

Ran 1 shell command

● Agent "Security review of deployed state" finished · 1m 5s

● **The gate passed: 0 blocking findings, 0 open questions, 1 info finding.**

The security, rules and drift reviewers all ran as ff-demo-reviewer.

Findings:
- **Info · SEC-010 · baseline · workspace FF Ticino Session 2 DEV**: The reviewer cannot read the workspace role assignments.
  - **Evidence:** GET roleAssignments returned HTTP 403 InsufficientWorkspaceRole.
  - **Remediation:** RLS (row-level security) does not apply to Admin, Member or Contributor, so a human needs to check who holds those roles.

Gate report: reports/20261028-080058-gate.md

I stopped at the gate. The snapshot is signed with ./py gate/snapshot.py sign only after you explicitly say "approve".

✻ Baked for 1m 18s · done 8:01 AM

›

›› auto mode on (shift+tab to cycle) · /tasks to see subagents

④ **Takeaway**

The rule you decided in step 3 also guards the next change. Rename it back, and the review passes.

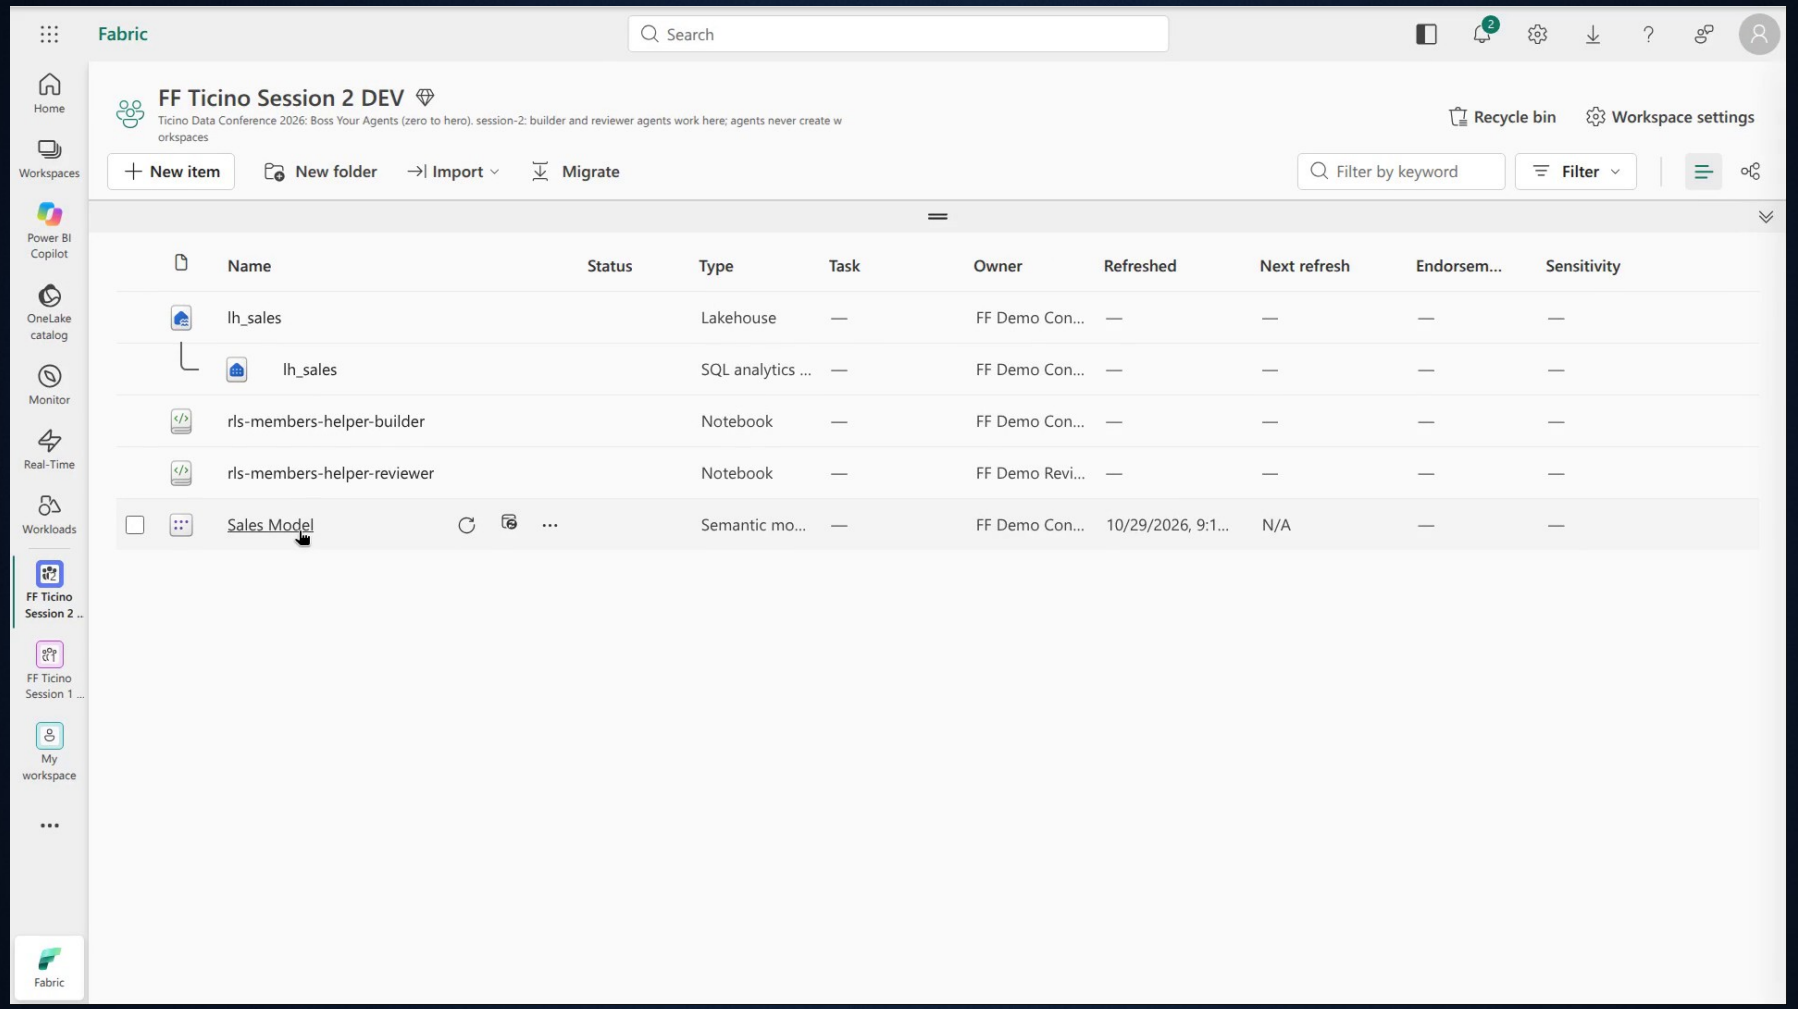
Fabric
Search
Home
Workspaces
Power BI Copilot
OneLake catalog
Monitor
Real-Time
Workloads
FF Ticino Session 2 ...
FF Ticino Session 1 ...
My workspace
Fabric
FF Ticino Session 2 DEV
Ticino Data Conference 2026: Boss Your Agents (zero to hero). session-2: builder and reviewer agents work here; agents never create workspaces
Recycle bin
Workspace settings
New item
New folder
Import
Migrate
Filter by keyword
Filter
Name
Status
Type
Task
Owner
Refreshed
Next refresh
Endorsem...
Sensitivity
lh_sales
Lakehouse
FF Demo Con...
lh_sales
SQL analytics ...
FF Demo Con...
rls-members-helper-builder
Notebook
FF Demo Con...
rls-members-helper-reviewer
Notebook
FF Demo Revi...
Sales Model
Semantic mo...
FF Demo Con...
10/29/2026, 9:1...
N/A

What the gate blocks

BLOCKS	REMEDIATION
SECURITY No RLS role on a table with sensitive data	Add a role that filters its rows, deploy, add members
SECURITY A role without members	Add the members; until then viewers see no data
SECURITY A role member reads other users' rows through another engine	Apply the same rule in every engine
DECISIONS The deployed state contradicts a recorded decision	Change the state back, or decide again and record it
RULES An accepted rule is violated	Follow the rule, or decide a new one and recompile the wiki
DRIFT Changed since the signed baseline	Undo the change, or review it and sign again

The gate stops the sign-off and the promotion, not the build in the DEV workspace.

FOUR PATTERNS

Where the four patterns sit in the steps

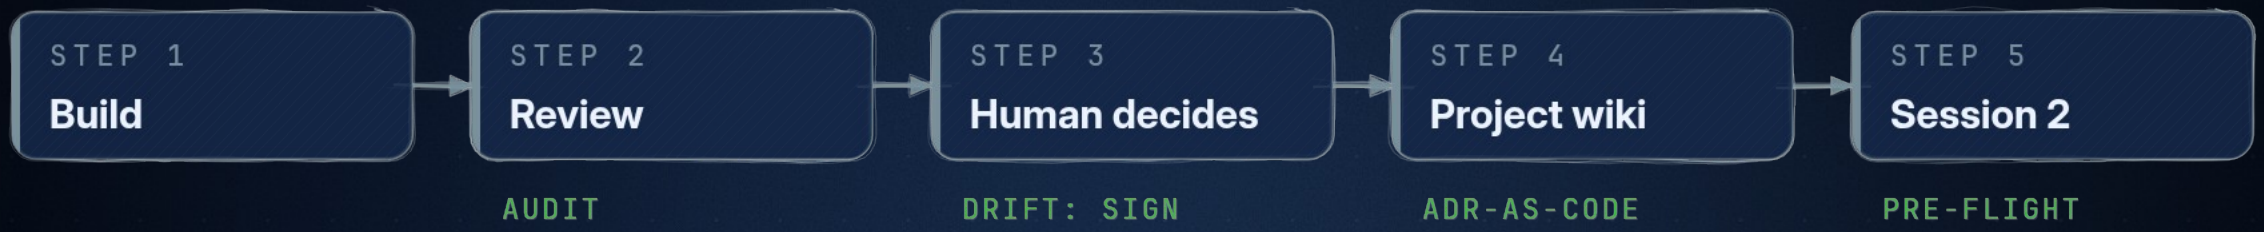

BOTH WORKSPACES ARE DEV · THE GATE STANDS BETWEEN DEV AND PRODUCTION

STEP 2

Post-deployment audit

Reviewers read what is deployed and test the effect: it caught a real RLS bypass.

STEP 3 · MONDAY

Drift detection

Signed at step 3, checked against after any change: the twist shows it.

STEP 4

ADR-as-code

A decision in the wiki, written so a reviewer can check it.

STEP 5

Pre-flight review

Session 2 plans against the wiki and builds both rules from the start: passed in round one.

AFTER THE RUN · WHAT STAYS

The project's knowledge, as plain files

THE FILES OF ONE RUN

```
session-1/
  CLAUDE.md   "Record every decision … as its own wiki page"
  .claude/commands/  build · decide · remediate · compile-wiki .md
review/.claude/agents/
  security-reviewer.md · decision-reviewer.md
  drift-reviewer.md
wiki/   17 pages + rules/
  index.md     one line per page
  ADR-005      Item naming with type prefixes
  ADR-006      Account-manager rule in every engine (model and SQL endpoint)
  SEC-020      RLS bypass through the SQL analytics endpoint
  DRF-000      No baseline yet; signed baseline of 2026-10-28
  rules/       ADR-001, -002, -004, -005, -006 .yaml
  8 pages      decision log: build, two decisions, the fix
session-2/
  CLAUDE.md   + "@wiki/index.md - read it and the pages it lists …"
```

ONE DECISION PAGE, VERBATIM

Direct Lake on OneLake (not on SQL endpoint)

DECISION

The semantic model Sales uses Direct Lake on OneLake: … partitions mode: directLake.

WHY

Microsoft recommends Direct Lake on OneLake for new semantic models; it never falls back to DirectQuery, and semantic-model RLS is supported.

ALTERNATIVES

Direct Lake on SQL analytics endpoint (can silently fall back to DirectQuery); Import (requires data refresh copies, not what was asked).

SOURCE

`learn.microsoft.com/fabric/fundamentals/direct-lake-how-it-works#directquery-fallback`

TAKE HOME

For Monday, in three sizes

TALK REPO

1. 30 minutes: CLAUDE.md instruction, Microsoft Learn MCP, a wiki folder
2. A day: reviewers with their own identity, a test identity in the RLS role
3. A sprint: rules, signed baseline, the next session with the wiki

Template in the repo · run it in a sandbox tenant first

SENTENCE OF THE DAY

Agents make good decisions when you let them record, check and hand them on.

You decide what becomes a rule.

Ticino Data Conference 2026

Thank you to our sponsors

Hosted by

acer

Gold

codeway

Lodestar

IT'S TIME TO INNOVATE

VARGROUP

Silver

Bi Factory

DATA KNOWLEDGE ADVISOR

CLOUDMETRIX

EY

K4ANALYTICS®

PLAN ANYWHERE

OKVIZ®

Thank you

Questions now, or later on LinkedIn

Matthias Falland

linkedin.com/in/matthias-falland

LINKEDIN

TALK REPO